

How the Public Defines Terrorism

Connor Huff Harvard University
Joshua D. Kertzer Harvard University

Abstract: Every time a major violent act takes place in the United States, a public debate erupts as to whether it should be considered terrorism. Political scientists have offered a variety of conceptual frameworks, but have neglected to explore how ordinary citizens understand terrorism, despite the central role the public plays in our understanding of the relationship between terrorism and government action in the wake of violence. We synthesize components of both scholarly definitions and public debates to formulate predictions for how various attributes of incidents affect the likelihood they are perceived as terrorism. Combining a conjoint experiment with machine learning techniques and automated content analysis of media coverage, we show the importance not only of the type and severity of violence, but also the attributed motivation for the incident and social categorization of the actor. The findings demonstrate how the language used to describe violent incidents, for which the media has considerable latitude, affects the likelihood the public classifies incidents as terrorism.

Replication Materials: The data, code, and any additional materials required to replicate all analyses in this article are available on the *American Journal of Political Science* Dataverse within the Harvard Dataverse Network, at <https://doi.org/10.7910/DVN/LD6NL8>.

Following the mass shooting in Charleston, South Carolina, there was widespread debate throughout the United States about whether to classify the violent incident as terrorism. FBI Director James Comey offered a negative assessment, noting that “terrorism is [an] act of violence . . . to try to influence a public body or citizenry, so it’s more of a political act. . . . [B]ased on what I know so far I don’t see [Charleston] as a political act” (Hattem 2015). His assessment was condemned from across the political spectrum, from critics on both the left (including then Democratic presidential candidate Hillary Clinton) and right (including then GOP presidential candidate Rick Santorum). Similar public debates erupted following other violent incidents, including the bombing at the Boston Marathon and shooting in Orlando, Florida. These debates highlight not only the contentiousness of classifying terrorism, but also the stakes involved in doing so, for policy makers, academics, and members of the public alike.

In this article, we turn to experimental methods to explain the tenor of these public debates. We investigate terrorism in a public opinion context not because we believe that the mass public’s intuitions can necessarily resolve normative debates about what should or should not be considered terrorism, but rather because of the central role that public opinion plays in our understanding of how terrorism works. In a vast array of prior research, terrorism is understood as a form of violence that functions by attracting public attention. It is because terrorism hinges on public reaction that Margaret Thatcher suggested terrorists depend on “the oxygen of publicity,” that Carlo Pisacane declared terrorism to be “propaganda by deed,” and that Ayman al-Zawahiri suggested that for al-Qaeda, media coverage is “more than half” the battle (Smith and Walsh 2013, 312). If the responses of ordinary citizens constitute a central causal mechanism through which terrorism operates, it logically follows that understanding *what ordinary citizens think terrorism is* is a crucial prerequisite to understanding *how they react to it*.

Connor Huff is Ph.D. Candidate, Department of Government, Harvard University, 1737 Cambridge Street, Cambridge, MA 02138 (cdezzanilhuff@fas.harvard.edu). Joshua D. Kertzer is Assistant Professor, Department of Government, Harvard University, 1737 Cambridge Street, K206, Cambridge, MA 02138 (jkertzer@gov.harvard.edu).

We thank Bob Bates, Mia Bloom, Alex Braithwaite, Dara Kay Cohen, Naoki Egami, Marcus Holmes, Jeff Kaplow, In Song Kim, Melia Pfannenstiel, Peter Vining, Ariel White, Thomas Zeitzoff, and audiences at the Political Violence Workshop at Harvard University, ISA 2016, MPSA 2016, Trinity College, Dublin, and William and Mary for helpful comments, suggestions, and assistance. Kertzer acknowledges the support of the Weatherhead Center for International Affairs, Institute for Quantitative Social Science, and Niehaus Center for Globalization and Governance; thanks to Perry Abdulkadir, Michael Chen, and Aaron Miller for fantastic research assistance. Authors’ names are in alphabetical order.

American Journal of Political Science, Vol. 62, No. 1, January 2018, Pp. 55–71

©2017, Midwest Political Science Association

DOI: 10.1111/ajps.12329

Employing a conjoint experiment on 1,400 American adults, we show how ordinary citizens classify terrorism based not only on relatively objective facts on the ground, but also on fairly subjective considerations about the perpetrator. On the one hand, considerations about the type and severity of violence matter—though interestingly, not the distinction between civilian and military targets that plays a central role in contemporary legal definitions. On the other hand, our respondents are also heavily influenced by descriptions about the perpetrator's identity and motivations, considerations about which the media has considerable latitude in the language it uses to cover incidents. In this sense, because actions do not always speak louder than words, the media has considerable power based on how it chooses to frame incidents: using predictive models derived from our experimental results shows how the likelihood Americans will classify incidents with the characteristics of the San Bernadino attacks, for example, can vary from around 31% to 82%, depending on the narrative constructed about the perpetrators' identity and motivations. In so doing, we empirically demonstrate the ways in which terrorism can be socially constructed.

The discussion that follows has four parts. First, we discuss the important role that public opinion plays in our understanding of how terrorism works, and we present a simple typology of factors that either loom large in formal terrorism classification schemes or in popular debates. Second, we discuss the design of our experiment. Third, we present our main results, whose implications we highlight using machine learning techniques to construct counterfactual simulations of real-world incidents, and automated content analyses of media coverage. We conclude by discussing the implications of our findings for the study of terrorism more generally.

The Stakes of Defining Terrorism

Whenever a major violent action takes place in the United States, a public debate erupts as to whether it should be classified as terrorism or not, mirroring a similarly contentious debate in the study of political violence as to how terrorism should be defined in the first place.

The contours of these debates matter for three reasons. First, classifying actions as terrorism has direct policy implications for how the perpetrators are prosecuted: in the United States, for example, terrorism is a federal charge, prosecuted in federal courts, whereas most violent crime is handled by the states. Similarly, transnational organizations deemed to be terrorist entities are subjected to harsh financial sanctions, foreign nationals on

terrorist watch lists are not allowed to enter the country, and Americans are prohibited from dealing with them. The consequences of acts being classified as terrorism are real and immediate. More generally, defining terrorism is crucial to fighting it: the United Nations has been unable to move forward with a comprehensive treaty against terrorism precisely because of the inability of its member states to come to a mutually acceptable definition of what terrorism is in the first place.

Second are important normative implications. As a discursive category, terrorism is understood as qualitatively different from other types of acts of violence: an extranormal or exceptional act, mandating an exceptional response (Agamben 2005). In the 2004 election campaign, when John Kerry pronounced that terrorism was ultimately a law enforcement issue, the George W. Bush campaign lampooned him in television advertisements for failing to take the threat seriously (Kertzer 2007, 964–65). Public debates following the shooting in Orlando in June 2016 displayed the same tendency. To categorize something as terrorism is to delegitimize its goals; terrorism is not merely a problem to be *managed*, but one to be destroyed; terrorists are to be *hunted*, rather than negotiated with (Hodges 2011). These normative questions are further exacerbated by complaints about the double standards with which the terrorist label is applied. When a group of armed white ranchers seized the headquarters of a U.S. Fish and Wildlife Service complex in Oregon in January 2016 to protest the federal government's policies on grazing and land rights, left-leaning commentators condemned the refusal of the media to label the group as terrorists, and bemoaned the seeming leniency the group was being given by law enforcement: "If Muslims had seized a federal building, they'd all be dead by now," the commentator Cenk Uygur wrote, and hashtags like #YallQaeda, #YokelHaram, and #VanillaISIS circulated on Twitter (Uygur 2016). Scholars have similarly noticed a double standard, in which the media is more likely to adopt an Islamic terror frame when the perpetrator is Muslim, and more likely to explore the attacker's personal life and mental health if the perpetrator is not (Powell 2011).

Third, these debates also have important stakes for political scientists. Almost every book or review article on terrorism we are aware of includes a compulsory line acknowledging the contentiousness of its definitional politics, and there are a variety of scholarly classification schemes that political scientists regularly employ.¹

¹ See, for example, Crenshaw (2000), Pape (2003), Horgan (2004), Goodwin (2006), Hoffman (2006), Kydd and Walter (2006), Merolla and Zechmeister (2009), Enders and Sandler (2011),

Indeed, a number of political scientists have used these typologies to constructively wade into public debates (Phillips 2015; Young 2012). Just as our understanding of the dynamics of civil wars changes when we adopt different operational definitions of the construct, our rules of thumb for how we know terrorism when we see it have important implications for which actions get turned into rows in our data sets and, thus, our fundamental understanding of how terrorism works.

In this article, we shift the focus from how governments, the media, and academics classify terrorism to instead explore how members of the mass public understand the term, as part of a growing body of research throughout the social sciences, using experimental methods to unpack our folk intuitions about political concepts.² We investigate terrorism in a public opinion context not because we believe that folk intuitions can resolve normative debates about what should and should not be considered terrorism—although disconnects between the judgments of citizens and the laws that govern them are indeed worthy of study (Audi 2009)—but rather because of the central role that public opinion plays in our understanding of how terrorism works. Most of our academic models of terrorism emphasize “victim-target differentiation,” or the notion that terrorism is a communicative act (“propaganda by deed”) directed at broader audiences (Asal and Hoffman 2015; Braithwaite 2013; Fortna 2015). In Bueno de Mesquita and Dickson (2007), terrorists carry out attacks in order to mobilize public support among the population they claim to represent; in Enders and Sandler (2011), terrorists seek to change government policy through public pressure; and in Kydd and Walter’s (2006) model of terrorism as costly signaling, three out of the five terrorist strategies they describe are directed at persuading the public. The public’s reactions are frequently posited to explain why terrorists adopt tactics like suicide bombing that tend to receive more publicity (Bloom 2005; O’Rourke 2009), and why democracies are more likely to be targeted by terrorist attacks (Eubank and Weinberg 1994; Pape 2003; Stanton 2013).

Unlike many foreign policy issues, terrorism is highly salient, capturing the public imagination—and producing downstream political consequences—to an extent that more ubiquitous forms of violence do not.³ It receives ex-

tensive media coverage (Nacos, Bloch-Elkon, and Shapiro 2011) and fuels powerful emotional responses that not only affect the public’s attitudes toward foreign policy issues (Albertson and Gadarian 2015; Huddy et al. 2005), but also how citizens act towards each other (Merolla and Zechmeister 2009). Politicians are aware of terrorism’s sway on the popular imagination, which is one reason why Democratic lawmakers use terrorism frames to pressure their Republican counterparts into supporting gun control measures (Cowan and Cornwell 2016). Outside the United States, political scientists have found that terrorist attacks increase support for right-wing parties in the locale where the attacks took place (Berrebi and Klor 2008), diminish willingness to negotiate (Huff and Kruszewska 2016), and sway elections (Bali 2007). Even the mere threat of terrorism is enough to change voting behavior (Getmansky and Zeitzoff 2014). Given this wealth of prior research in which public reactions play an important role as a causal mechanism linking violent incidents to downstream political consequences, it is crucial for political scientists to understand what the public thinks it is reacting to.

A Typology for Classifying Terrorism

In this section, we present a typology for classifying terrorism, integrating common components of formal classification schemes with the contentious elements of public debates that erupt in the wake of violent incidents. The typology is composed of two broad components. The first consists of relatively objective facts on the ground: information about the type and severity of the violence employed, and the incident’s target and location. This information is often available immediately after the incident occurs. In contrast, the second consists of information pertaining not to “what” questions, but “who” and “why,” concerning the identity of the perpetrators and motivation attributed to them. This category thus consists of considerations that are often relatively subjective, are unavailable until days or weeks after an incident occurs, and give the media more leeway in how they frame events. We discuss each element in turn.

Objective Criteria: The Facts on the Ground

The first component focuses on the fundamental descriptive characteristics of the action. Scholarly typologies and

merely “one among many important issues.” In contrast, the same poll found that although 74% of Democrats called mass shootings a critical issue, only 60% of Republicans did (Cox and Jones 2015).

Phillips (2014), Fortna (2015), Horowitz (2015), Huff and Kruszewska (2016).

²See, for example, Malle and Knobe (1997), Kertzer and McGraw (2012), and Kertzer, Renshon, and Yarhi-Milo (2016).

³For example, a December 2015 Public Religion Research Institute poll found that 84% of Republicans and 70% of Democrats saw terrorism as a “critical issue” facing the country, as opposed to

definitions of terrorism consistently emphasize two types of factors when discussing terrorist incidents. The first is *violence*, mentions of which appear in almost every legal and academic definition alike (Weinberg, Pedahzur, and Hirsch-Hoefler 2004). In the discussion below, we suggest violent incidents can be understood along two key dimensions: the type of tactic used, and the severity of the consequences as measured by casualties. The second component is information about whom the violence was against, conveyed through either the type of target or location of the incident. We discuss each component in turn.

Type and Severity of Violence. We argue that the violence observed in incidents can be understood along two dimensions. First, we can think of different “technologies” (Kalyvas and Balcels 2010), or “repertoires of contention” (Tilly 1986), characterized by the particular tactics employed in an incident. For example, bombing is a different type of violence than shooting or hostage taking. Following the Boston Marathon Bombing in 2013, President Barack Obama stated that “any time bombs are used to target innocent civilians, it is an act of terror” (Landler 2013). This statement is emblematic of the close mapping in both academic and policy circles between the type of violence used and the likelihood the incident will be classified as terrorism, buttressed by recent work in political science on the role that different types of violent tactics play in structuring perceptions of actors (Abrahms 2013; Huff and Kruszezewska 2016). The empirical implications of this prior research are twofold. First, violent tactics should be more likely to be classified as terrorism than nonviolent tactics. Second, bombings are a unique form of violence often associated with terrorism (Bloom 2005), such that we expect them to be more likely to be classified as terrorism than other forms of violence.⁴

In addition to characterizing violence by its type, it can also be characterized by its *severity*. Indeed, the magnitude of the carnage imposed by an attack is often the central focus of the media following violent incidents, with headlines typically bringing the number of casualties to the fore. This leads to a second empirical implication: the higher the number of casualties, the more likely the incident will be classified as terrorism.

The Target and Location. We can characterize the target of violent incidents in two ways. First is whether the target is a civilian or noncombatant (Coady 2004; Rodin 2004). Consistent with the principle of distinction

between combatants and noncombatants in *jus in bello*, the logic is that attacks upon the government or state apparatus might be undesirable, but they are nonetheless legitimate in a way that targeting civilians is not. The clear empirical implication is that attacks on targets closely associated with formal state institutions are more likely to be considered legitimate, and thus should be less likely to be considered terrorism.

Second is the broader location in which the target is situated. This distinction was salient following the terrorist attacks in Paris and Lebanon in the fall of 2015, when a number of pundits explicitly criticized Western publics for the selective empathy revealed by their willingness to change their Facebook profile pictures in solidarity with Paris while ignoring the attacks that had occurred in Lebanon the previous day. Consistent with research on ingroup favoritism (Brewer 1999) and affective responses decreasing with social and spatial distance (Lieberman, Trope, and Stephan 2007, 373), the empirical implication is that Americans should be more likely to classify events as terrorism when the incident takes place in the United States, compared to in foreign countries, particularly for foreign countries unlike the United States.

Subjective Criteria: The Categorization of the Actor and Motivation

If the previous definitional components focus on the *act* carried out, another set emphasizes the *actors* themselves and the *motivation* attributed to their behavior. This idea is most commonly embedded in formal definitions noting that actors must have engaged in the act for “political aims and motives” (Hoffman 2006), “achieving a political goal” (Keller 2005), or “the advancement of some political, ideological, social, economic, religious, or military agenda” (Shanahan 2010).

Three points are worth noting here. First, determining motivations is difficult, since it requires overcoming “the problem of other minds” and accessing an interior attribute of other actors, who may either have strategic reasons to misrepresent, or who themselves may not be aware of their underlying reasons for action. It is for this reason that classic international relations (IR) scholars like Morgenthau (1985, 5) famously discouraged political scientists from studying motivation altogether. At the same time, however, as political scientists (Herrmann 1988) and psychologists (Malle and Knobe 1997) note, we intuitively rely on assumptions about motivations both when explaining actions and evaluating the actors who carried them out.⁵ In other words, the difficulty of

⁴This could either be due to the frequency with which the public views bombings as being conducted by organizations associated with terrorism, or due to their often indiscriminate nature (Goodwin 2006).

⁵See, for example, work in IR on the security dilemma.

assessing motivations does not deter us from relying on them.

Second, as a result, we often infer motivations for action using information about the actors themselves; we tend to provide answers to questions about “why” using answers to questions about “who.” It is thus worthwhile to parse out which particular components of descriptions of actors lead individuals to classify acts as terrorism. Third, precisely because of the subjectivity inherent in interpreting actions, the media has a considerable amount of leeway in the narrative it constructs about the actor, from the characteristics it mentions to the rationales it attributes. We discuss each of these considerations in turn.

Type of the Actor. Following violent incidents, governments, the media, and the public search for who is responsible, whether individuals or collectives. The differences are embodied in the distinction between “lone wolves” and “terrorist organizations.” We argue that we should be more likely to expect incidents perpetrated by organizations to be classified as terrorism than those by individuals. Because we tend to think of terrorism as “a collective, organized activity” (Spaaij 2010, 855) and expect individuals to infer that formal organizations act strategically in pursuit of some broader political or social agenda, it follows that incidents carried out by lone individuals are less likely to be seen as terrorism than incidents carried out by groups, particularly more formal organizations. This means that when a bombing or hostage taking is carried out by an organization, for example, we tend to be more likely to assume it was carried out in pursuit of a broader political agenda than if the same incident were carried out by a “lone wolf.”

By contrast, in the wake of violent incidents perpetrated by individuals, there is far more ambiguity about whether the action was taken purposively in pursuit of a broader agenda. Given the kaleidoscope of violence undertaken by individuals in the United States, knowing that an action was undertaken by an individual provides less information about whether an incident was undertaken for a broader political purpose than whether an incident was perpetrated by an organization. In this sense, providing information about the type of actor provides information about why the event occurred.

We can use similar logic to consider incidents perpetrated by individuals with a history of mental illness. These types of incidents stir some of the most controversial public debates on when and how events should be defined as terrorism. For example, following recent shootings in Fort Hood, Charleston, and Orlando, pundits debated whether the perpetrators’ history of mental illness should affect how we judge their actions. Just as

we expect incidents perpetrated by individuals to be less likely to be classified as terrorism than those carried out by more formal organizations, we can similarly expect incidents perpetrated by individuals with a history of mental illness to be less likely to be attributed to a purposive political agenda, and thus less likely to be considered terrorism.

Actor Categorization and Motivation. The broader motivations for incidents are often understood in two ways. The first is through the identity categories used to describe the type of actor. Because of the pervasiveness of social categorization processes (Brewer 1999), these descriptions both provide information about the actor and enable observers to draw inferences about broader political agendas. Some have argued that what matters is not only that the actor is seen as fighting for a broader political purpose, but also the particular agenda for which the actor is mobilized, hence the flurry of debates about the presence of “double standards” in whom we define as terrorists. For example, it might matter if the actor is described as Muslim, rather than Christian, or left-wing, rather than right-wing. Most broadly, we hypothesize that incidents perpetrated by an actor described as Muslim are more likely to be classified as terrorism than other types of descriptions (D’Orazio and Salehyan 2017; Lemieux and Karampelas 2016). The second is through the explicit motivation attributed to the action. For example, was the actor motivated by hatred toward the target, a personal agenda, or a desire to change government policy? Each of these types of motivations provides varying levels of information about the purposiveness of the action and, thus, the likelihood we should expect an incident will be perceived as terrorism.

Method

In this article, we seek to directly assess the attributes of incidents ordinary citizens use to define incidents as terrorism. We do so using a conjoint experimental design (Hainmueller, Hopkins, and Yamamoto 2014; Kertzer, Renshon, and Yarhi-Milo 2016) in which we present our participants with a series of incidents with randomly generated features and then ask whether they would classify each incident as terrorism. The use of a ratings-based conjoint design offers two main advantages. First, we can hold fixed a range of attributes of incidents that could confound observational studies on the topic. Even if there were an abundance of high-quality public opinion data on how the public defines terrorism, it would be extremely

difficult to disentangle the effects of our attributes of interest, particularly given the presence of both incident-specific and time-varying characteristics of the political environment surrounding the event, as well as our interest in the different types of media frames that prior research has shown vary with the type of actor. Given these problems, experiments give researchers vital control. Second, the use of a conjoint experimental design allows us to study the effect of each attribute, using a relatively large number of experimental treatments, in a manner that would not be possible with more traditional factorial experiments due to statistical power constraints. While prior experimental research has sought to explore the effects of some of the factors we discuss in this article, a conjoint design allows us to explore how they operate in tandem.⁶

The Survey Instrument

In the experiment, we manipulate seven attributes of an incident. As noted above, each of the attributes of the incident chosen, and the treatments within them, were included to reflect long-standing debates within both academic and policy worlds about how we should define terrorism. Each of these attributes is manipulated within a text block where the structure is intended to mirror what respondents would observe in the first few sentences of a newspaper article describing a recent incident. After respondents read about the details of the incident, they were then asked whether they would classify it as terrorism. Respondents went through this process seven times in total. Each of these attributes is presented in Table 1 and discussed in the text below.

First, we manipulated the *tactics*: whether the actor carried out a protest, hostage taking, shooting, or bombing. Protests were chosen as a baseline category of nonviolent political action to test how the use of more violent tactics affects whether individuals perceive an incident to be terrorism. Shooting and hostage taking were chosen to represent what we generally think of as violent tactics that may or may not be associated with formal organizations or terrorists. Finally, bombing represents one of the most emblematic tactics associated with terrorism. Each of these tactics also has the benefit of plausibly being adopted by a range of actors with varying ideologies, targets, and motivations.

Second, we manipulated the severity of the *casualties*, which was set to vary from either none, one, two, or

10. We specified these casualty levels with three goals in mind. First, they serve as a way of unpacking the marginal effect of individuals dying as the result of an incident, to capture the notion that there is something systematically different about incidents that result in fatalities from those that do not. Second, they explore the functional form of the effect: is a single casualty enough for an incident to be seen as terrorism, or does the likelihood increase with severity? Finally, the casualty levels are calibrated to reflect the real-world distribution of casualty levels for incidents commonly classified as terrorism, while also ensuring sufficient variation to be able to reflect higher-casualty incidents.⁷

Third, we manipulated the *target* of the incident, which was either a military facility, a police station, a school, a Christian community center, a Muslim community center, a Jewish community center, a church, a mosque, or a synagogue. The use of a military facility allows us to have a target not associated with civilians, and police stations reflect targets that are affiliated with the government but generally not perceived to be engaged in combat. The remaining seven categories represent different types of civilian targets, with the use of a school capturing a type of target that is affiliated directly with neither combat activities nor religion. Finally, we distinguish between religious places of worship and religiously affiliated community centers, due to the heightened symbolism of attacking the former.

Fourth, we manipulated the *location* of the incident, such that the act took place in either the United States, a foreign democracy, a foreign democracy with a history of human rights violations, a foreign dictatorship, or a foreign dictatorship with a history of human rights violations. We varied the location to explicitly test how differences in the attributes of the target population and government affect whether incidents should be perceived to be terrorism. By manipulating the location of the incident, we are able to explicitly test whether and how incidents more closely linked to Western institutions affect whether respondents are more likely to classify an event to be terrorism.

Fifth, we manipulated the *actor* who carried out the incident, varying whether the perpetrator was an organization, an organization with ties to the United States, an organization with ties to a foreign government, a group, an individual, or an individual with a history of mental illness. Employing these different types of perpetrators

⁶See Appendix S1.2 in the supporting information (SI) for further discussion of ratings-based conjoints.

⁷These figures reflect the low number of individuals typically killed in these incidents. In the Global Terrorism Database (GTD), for example, the median bombing has no fatalities, a bombing that kills 10 people is in the 96th percentile, and a hostage taking that kills 10 people is in the 98th percentile.

TABLE 1 Conjoint Treatments

(A) <i>Tactic</i>	The ... (1) ... protest (2) ... hostage taking (3) ... shooting (4) ... bombing
(B) <i>Target</i>	... occurred at a ... (1) ... military facility (2) ... police station (3) ... school (4) ... Christian community center (5) ... Muslim community center (6) ... Jewish community center (7) ... church (8) ... mosque (9) ... synagogue
(C) <i>Location</i>	... in ... (1) ... the United States. (2) ... a foreign democracy. (3) ... a foreign democracy with a history of human rights violations. (4) ... a foreign dictatorship. (5) ... a foreign dictatorship with a history of human rights violations.
(D) <i>Casualties</i>	There ... (1) ... were no individuals (2) ... was one individual (3) ... were two individuals (4) ... were ten individuals ... killed in the [tactic].
(E) <i>Actor Description</i>	The [tactic] was carried out by ... (1) ... an (2) ... a Christian (3) ... a Muslim (4) ... a left-wing (5) ... a right-wing
(F) <i>Actor Type</i>	... (1) ... organization. (2) ... organization with ties to the United States. (3) ... organization with ties to a foreign government. (4) ... group. (5) ... individual. (6) ... individual with a history of mental illness.
(G) <i>Actor Motivation</i>	News reports suggest ... (1) ... that there was no clear motivation for the incident. (2) ... the incident was motivated by the goal of overthrowing the government. (3) ... the incident was motivated by the goal of changing government policy. (4) ... the incident was motivated by hatred towards the target. (5) ... the individual had been in an ongoing personal dispute with one of the targets.

allows us to test the ways in which the perceived “purposiveness” of the actor affects the likelihood an incident is perceived to be terrorism.

Sixth, we manipulated the social *categorization* of the perpetrator, either describing the actor as Christian, Muslim, left-wing, or right-wing, or assigning the actor to a control condition in which no categorization was given, acting as a baseline from which to interpret the effects of the other descriptions. We chose these categories as they loomed the largest in recent prominent debates about classifying incidents as terrorism.

Finally, we manipulated the *motivation* attributed to the perpetrator: government overthrow, policy change, hatred toward the target, a personal dispute with the target, or no clear motivation. By including a category where the motivation is unclear, we can mirror the general uncertainty and absence of public information immediately following a violent incident, approximating the actual dissemination of information as an incident unfolds. The government overthrow and policy change motivations allow us to capture some of the most prominent goals of terrorist organizations as studied by political scientists (Kydd and Walter 2006), whereas attributing the incident to hatred toward the target allows us to capture some of the more contentious incidents subject to recent debates. Finally, the use of a personal dispute as a baseline category allows us to study the effects of political motivations versus apolitical ones.

To be clear, our claim is not that the seven sets of factors discussed above and summarized in Table 1 constitute a complete set of characteristics that could be relevant to debates about terrorism, but rather that they mirror a wide set of characteristics from either public debates or the academic literature. Indeed, despite not being an exhaustive list, we note that a fully crossed factorial featuring all permutations of the factors we manipulate here would contain 108,000 combinations.⁸ An example of how the vignette appeared to respondents is presented below in Figure 1.

Fielding the Experiment

The experiment was embedded in a survey fielded on 1,400 adults recruited using Amazon’s Mechanical Turk

⁸Thankfully, as noted above, conjoint experiments leverage between- and within-subject variation across participants to achieve higher levels of statistical power. Moreover, following best practices, we included randomization constraints to prevent implausible or problematic combinations of incident attributes from skewing the results. See Appendix §1.2 in the SI.

(MTurk) in August 2015.⁹ One commonly voiced concern with MTurk samples is that although they are relatively diverse, they are not representative of the American population as a whole (Huff and Tingley 2015). To mitigate against biased estimates of our treatment effects, for the main analyses presented below we employ entropy balancing to reweight the data to population parameters (Hainmueller 2012; for a similar empirical application, see Kertzer et al. 2014), though in Appendix §1.3 in the SI, we show that the substantive conclusions remain the same regardless of whether weights are used. Because we have 1,400 participants who classify seven incidents each, the analyses that follow are of 9,800 different randomly generated scenarios, analyzed with clustered, bootstrapped standard errors at the participant level.

Results

In this section, we present the experimental results.¹⁰ We do so in three parts. First, we demonstrate the importance of the type of tactic and number of casualties. In contrast, the target and location of the incident have no significant effect. Second, we show that the perceived political purposiveness of the actor, inferred either indirectly through their social categorization or directly through their posited motivation, has a significant effect on the likelihood an incident is classified as terrorism. In general, the perception that an actor is acting purposively for some broader political agenda matters more than the content of the agenda itself.

The Type of Tactic and Severity of the Violence


The results suggest that the *type* of violence plays a central role in determining whether an incident is classified as terrorism, whereas the *severity* of the violence seems to be much less central. The left-hand panel of Figure 2 shows the relative importance of the tactics chosen by the actor regarding whether an incident is classified to be terrorism. As is standard in conjoint experiments, our quantities of interest here are average marginal component effects (AMCEs), depicted in a probability scale on the x-axis, such that point estimates further to the right

⁹For a more detailed discussion of the timing of the experiment and how it might substantively affect results, see Appendix §1.1 in the SI. See Appendix §1.4 for a discussion of MTurk.

¹⁰See Appendices §1–4 in the SI for a range of additional methodological information and analyses.

FIGURE 1 An Example Vignette

The incident: a shooting.
 The shooting occurred at a church in a foreign democracy with a history of human rights violations.
 There were two individuals killed in the shooting.
 The shooting was carried out by a Muslim individual with a history of mental illness.
 News reports suggest the individual had been in an ongoing personal dispute with one of the targets.

FIGURE 2 The Effect of Tactics and Casualties

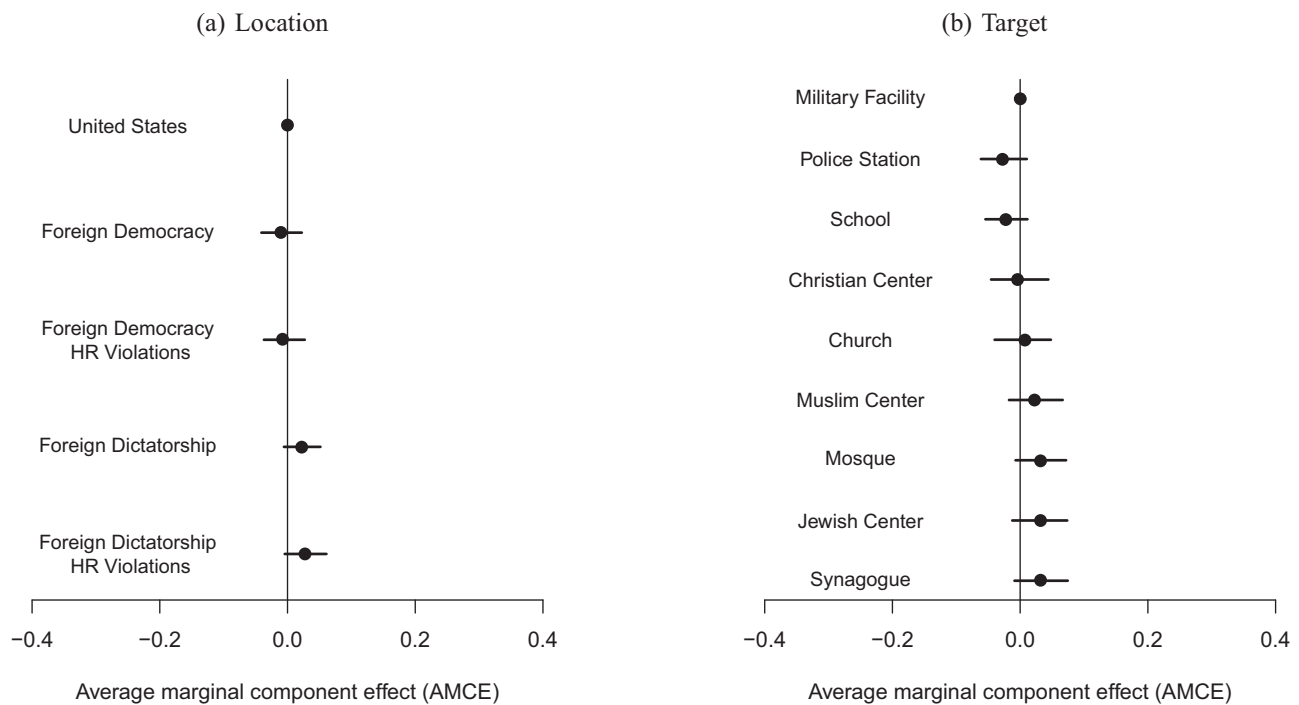
indicate a greater probability that an event will be classified as terrorism; as in the figures that follow, the point estimates are depicted with 95% confidence intervals derived from $B = 1,500$ clustered bootstraps. Two patterns are evident. First, the use of violence overwhelmingly increases the likelihood an incident will be classified as terrorism, consistent with the prominent role violence plays throughout scholarly definitions. Second, the type of violence matters as well, as bombings are significantly more likely than either shootings or hostage taking to be classified as terrorism. This finding is consistent with the statement made by President Obama in the wake of the Boston Marathon Bombing, which implied a uniqueness about bombing as a tactic overwhelmingly associated with terrorist organizations. The results presented here suggest this perception is shared by ordinary citizens.

The right-hand panel of Figure 2 shows the effect of increasing casualties on whether an incident is considered

to be terrorism. Although higher-casualty incidents are more likely considered to be terrorism, the size of the effect is small relative to the increase associated with violent tactics. In supplementary analyses in Appendix S2.2 in the SI, we reiterate the point by showing that a bombing with no fatalities is statistically indistinguishable from a shooting that kills 10 people.

The Target and Location

Figure 3 demonstrates that there is no significant relationship between either the target or location of an incident and the likelihood it is classified as terrorism. Whether the incident occurred at home or abroad, or in countries with different regime types or human rights records appears to be immaterial. We therefore find no evidence that respondents employ differential definitions in contexts

FIGURE 3 Location and Target of Incident


where the attack might be “justified” by prior government repression.

We observe similarly null effects for the target of the incident: there is no statistically significant difference between how respondents view incidents targeting military facilities and other government, civilian, or religious centers. This absence is striking given the frequency with which terrorism definitions often distinguish between civilian and military targets; ordinary citizens seem not to make this distinction.¹¹

Attributes of the Actor: Type, Descriptions, and Motivation

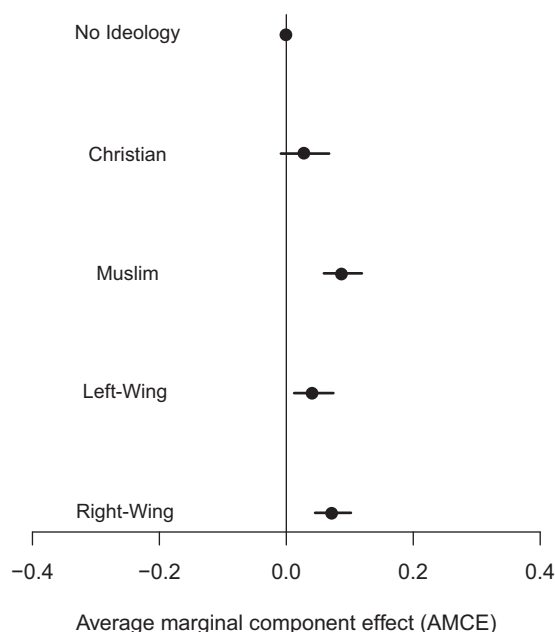
Thus far, we have shown that both the type and severity of the violent tactics employed strongly shape the likelihood individuals will classify an event as terrorism, whereas the target and location of the attack do not. The above factors are relatively objective “facts on the ground,” information available relatively shortly after an incident has taken place. In contrast, we now turn to relatively subjective characteristics about the actor itself about which the media has much greater latitude when calibrating coverage.

In Figure 4, we present the results from our actor treatment, which reveal two main patterns. Acts carried

FIGURE 4 The Political Purposiveness of the Perpetrator

out by collectives are around 15% more likely to be understood as terrorism than acts carried out by individuals. The effect is particularly strong for organizations with ties to foreign governments, showcasing the extent to which


¹¹Though see Young and Findley (2011, 299–300).

FIGURE 5 Social Categorization Effects

terrorism is thought of as politically purposive. Consistent with this emphasis on purposive action, incidents perpetrated by an individual with a history of mental illness are significantly less likely to be classified as terrorism.

What is especially interesting about these findings is that some of these considerations are relatively subjective, thereby highlighting the way in which variation in the presentation of information about the actor affects whether it is perceived as terrorism. This means that the media have the ability to shape perceptions of violent incidents, and subsequently whether they are classified as terrorism, in the way that they choose to frame events, a point we explore in detail below.

We see a less dramatic pattern in Figure 5, which depicts how the way in which the actor was described shapes perceptions of whether an incident should be classified as terrorism. The left-hand panel shows how providing a social categorization (particularly Muslim or right-wing) significantly increases the likelihood an incident will be classified as terrorism. The effect is strongest for perpetrators described as Muslim, though even then, the effect size remains relatively small compared to the other factors presented above. Interestingly, incidents where the perpetrator is described as Christian are not significantly more likely to be classified as terrorism. The fact that we find these effects in spite of potential social desirability biases suggests we should think of the magnitude of our Muslim effect as an underestimate, rather than an overestimate,

FIGURE 6 Motive Attribution Effects

although supplementary analyses reported in Appendix §2.3 in the SI suggest the magnitude of the bias is small.

Finally, Figure 6 shows how the motivation attributed to the incident affects whether an incident is perceived to be terrorism. The main findings are threefold. First, incidents motivated by a personal dispute with the target are the least likely to be understood as terrorism. This matters given recent debates over whether, for example, the shooting of three Muslim university students in Chapel Hill, North Carolina, in 2015 (ostensibly over a parking dispute) was terrorism or not. Second, compared to personal disputes, we find that an incident where the motivation is unclear increases the likelihood it is understood as terrorism: when in doubt, our respondents are more likely to assume an incident is political rather than personal. Third, we find that from a baseline where the motivation is unclear, the motivations of hatred, policy change, and government overthrow increase the likelihood an incident is classified as terrorism. The powerful role of hatred here (an effect not statistically significant from that of policy change) is especially interesting since it does not exert a similarly large presence in legal definitions of terrorism, which tend to emphasize formal policy goals rather than animosities. Nonetheless, the strong effects of policy-oriented goals here reinforce the extent to which our participants associate terrorism with acting purposively for a broader political agenda.

Extensions and Applications

The strength of the conjoint experiment is that it lets us examine the relative contributions of a large number of factors in shaping how ordinary citizens think about terrorism. However, it also has another virtue: we can invert our experimental findings to model how the public thinks about real-world incidents. Doing so has three benefits. First, it allows us to code a range of incidents that have occurred in the real world using the typology presented above in order to predict the probability that participants would classify each event as terrorism. Second, we exploit the subjectivity in many of these coding decisions to show how we can construct a number of different, equally plausible narratives about the same event, leading to sizable differences in the proportions of the public that should understand the event as terrorism or not. In this sense, we show empirically how terrorism can be socially constructed through media and elite framing. Finally, we can validate our findings using automated content analyses of media coverage. We sketch out all three, in turn.

Application I: Mapping onto Real-World Events

First, we code 39 recent or otherwise prominent incidents that can be subsumed by the typology presented above. Because of the potential nonadditive effects of the attributes we explore here, we employ a machine learning method proposed by Imai and Ratkovic (2013) and Egami and Imai (2015), who apply a variable selection problem approach to the study of high-dimensional treatment interactions by adapting a support vector machine classifier with LASSO constraints, thereby estimating a high-dimensional, interactive model without overfitting the data.¹² We then calculate fitted values to estimate the predicted probability that each incident would be understood as terrorism, presented in Table 2.

Several patterns in Table 2 are striking, including the extent to which bombings perpetrated by formal organizations (e.g., Hamas, Hezbollah, ETA, and the Ku Klux Klan) are the types of incidents most likely to be perceived as terrorism, as well as the proportion of incidents with predicted probabilities clustering around 50%, illustrating the contentiousness and difficulty of defining terrorism in the wake of violent incidents.

¹² Consistent with the main results presented throughout this article, we present weighted estimates here; see Appendix S4 in the SI for alternative modeling specifications.

Application II: Demonstrating Framing Effects

Especially important for our purposes, though, is that the ranking of many of these incidents is in part a function of relatively subjective coding decisions regarding the description of the actor and motivation. In our view, this is one of the most striking and important implications of this article, since the media make similar coding decisions when choosing whether to highlight a suspect's religion, speculate on foreign ties, attribute a political motivation behind the incident, or psychoanalyze the perpetrator from a distance.

To make this discussion concrete, consider the shooting that took place in December 2015 in San Bernardino, California. In the aftermath of the shooting, a number of facts emerged about the two perpetrators: Tashfeen Malik had pledged allegiance to ISIS in a Facebook post, and Syed Rizwan Farook's father was described as being both abusive and mentally ill. There are at least three subjective decisions one can make about how to code the event. First, was the act carried out by lone individuals, or were they actors with foreign ties? Is pledging allegiance on social media sufficient to constitute a foreign tie? Second, was the attack motivated by policy goals, or was the motivation behind the shooting unclear? Third, is shooting up an office party indicative of mental illness, or merely malign intentions? Our model suggests these coding decisions have important implications for how people understand the attack. If the perpetrators are coded as having foreign ties and the goal of changing policy, our model suggests an 82% likelihood of its being terrorism. If, instead, the perpetrators are simply seen as individuals with no clear motivation, the probability of classifying the incident as terrorism drops to 50%. If the relevant narrative instead raises the specter of mental illness and does not mention the attackers' religion, the probability drops to 31%. In this sense, there is over a 50% difference in the likelihood the attacks will be classified as terrorism, based solely on how the story is framed.

Application III: Analyzing Media Coverage

The results in Table 2 also suggest two additional implications. First, if our classifications are relatively accurate, they should correlate with the rate at which the incidents are described as terrorism in the real world: an event with the characteristics that our participants see as highly likely to be terrorism should generally be described as such in the press, for example. Second, the extremity of the probability estimates should be negatively correlated

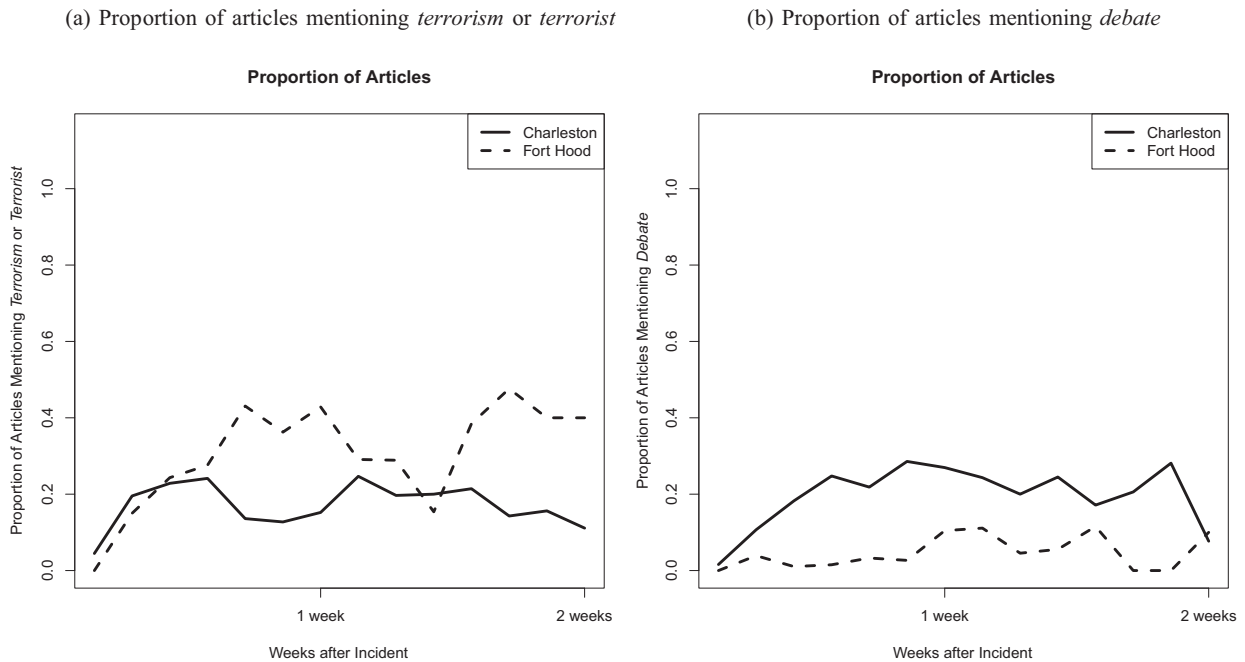
TABLE 2 The Predicted Probability a Range of Incidents Are Classified as Terrorism Using a Weighted Support Vector Machine Classifier with Lasso Constraints

Incident	Date	Predicted Probability
East Selma Church Shooting	09/20/15	0.23
Dallas Police HQ Shooting	06/13/15	0.25
UCLA Black Lives Matter Protest	10/08/15	0.28
University of California Tuition Hike Protests	03/18/15	0.31
Zvornik Police Station Shooting	04/27/15	0.33
Marysville Pilchuck High School Shooting	10/24/14	0.33
Rocori High School Shooting	09/24/03	0.33
Shooting of Police in Brooklyn	12/20/14	0.35
Islamic Community Center of Phoenix Demonstrations	10/10/15	0.35
Shooting of Police in Oakland	03/21/09	0.35
Pentagon Metro Shooting	03/04/10	0.36
St. Columbanus Church Shooting	11/26/12	0.39
Poe Elementary School Bombing	09/15/59	0.39
Newport Church Hostage Situation	07/30/06	0.39
Copenhagen Synagogue Shooting	02/14/15	0.39
University of Alabama Huntsville Shooting	02/12/10	0.43
Charleston Church Shooting	06/17/15	0.46
Camp Shelby Shootings	08/05/15	0.47
Overland Park Jewish Community Center Shooting	04/13/14	0.47
Knoxville Unitarian Universalist Church Shooting	07/27/08	0.48
Rosemary Anderson High School Shooting	12/13/14	0.49
Shooting of George Tiller	05/31/09	0.55
Bombing of Shiraa Village Mosque	12/30/14	0.55
Kehilat Bnei Torah Synagogue Attack	11/18/14	0.59
KKK Selma Bombing	09/15/63	0.61
Seattle Jewish Federation Shooting	07/28/06	0.62
Nag Hammadi Massacre	01/07/10	0.62
Contra Attack in Quilali	11/11/87	0.62
ETA Sanguesa Car Bombing	05/30/03	0.64
Lombard Islamic School Bombing	08/12/12	0.71
Hamas Attack on IDF in Khan Yunis	12/24/14	0.71
Shebaa Farms Incident	01/28/15	0.72
Camp Integrity Suicide Bombing	08/07/15	0.74
Porte de Vincennes Hostage Situation	01/09/15	0.74
Pakistan Army General HQ Hostage Situation	10/10/09	0.78
Zif School Bombing	09/17/02	0.84
Aksu Bombing	08/19/10	0.85
Chattanooga Shootings	07/16/15	0.85
Fort Hood Shootings	11/05/09	0.85

with the magnitude of public debate: events clustered at around 50% on our probability scale should be the ones that produce the greatest dissensus.

Automated content analyses of newspaper coverage thus serve as a helpful means of validating our findings. As a plausibility probe, we selected two events from

Table 2: an incident with a high classification rate (the shooting at Fort Hood, with a terrorism probability of 85%), and an incident with a classification rate close to 50% (the shooting at Charleston, with a terrorism probability of 46%). We then used LexisNexis to collect every newspaper article about the incidents published in the

FIGURE 7 Automated Content Analysis of Media Coverage

Note: Automated content analysis of daily newspaper coverage of the attacks in Charleston and Fort Hood (which our models classify as 86% and 46% likely to be classified as terrorism, respectively) provides results consistent with our theory. Panel (a) shows that newspaper coverage of Fort Hood is consistently more likely to refer to terrorism than articles covering Charleston. Panel (b) shows that newspaper coverage of Charleston was significantly more likely to discuss the presence of “debate” than newspaper coverage of Fort Hood.

two weeks following each attack.¹³ There are several important similarities between these incidents: for example, both shootings received high levels of media coverage, were perpetrated by an individual, and took place in the United States. Yet there are also a number of important differences that map directly onto the typology presented in this article: for example, the perpetrator in Charleston had a more pronounced history of mental illness, was associated with right-wing extremists, and was described as being motivated by hatred rather than formal policy change.

For each event, we calculated two quantities of interest. The first is the proportion of articles published making reference to the term *terrorism* or *terrorist*: if our model is accurate, we should expect that news coverage of the incident with a high predicted classification rate should more frequently invoke terrorism than the coverage of the event with a moderate predicted classification rate. The second is the proportion of articles published making reference to “debate”: if our model is reliable, we should expect that the coverage of the event

with a high predicted classification rate should feature less debate than the coverage of the event with a predicted classification rate that hovers near a coin toss. Sure enough, when we plot the daily counts of both quantities in Figure 7, we find that articles about the shooting at Fort Hood were more likely to mention the word *terrorism* or *terrorist* than articles about the shooting in Charleston. Indeed, in the two-week window around each of the respective incidents, articles about Charleston were approximately 13% less likely to mention the word *terrorism* or *terrorist* than articles about Fort Hood ($p < .01$), consistent with the idea that incidents we predict as more likely to be classified as terrorism are indeed discussed as such in the news media.

Similarly, the right panel of Figure 7 demonstrates that articles about the shooting at Charleston were more likely to mention the word *debate* than articles about the shooting at Fort Hood. In the two-week window around the incidents, articles about Charleston were 15% more likely to mention the word *debate* than articles about Fort Hood. Although more of a plausibility probe than a full test of our theory, these results show how media coverage offers another domain in which our theoretical framework can be explored.

¹³See Appendix §6 in the SI for details.

Conclusion

This article contributes to broader research on the relationship between public opinion and government action in the wake of violent events by exploring the characteristics that shape how ordinary citizens classify incidents as terrorism. A rich literature in IR explores the distinction between terrorism and “big” forms of violence like civil wars (e.g., Findley and Young 2012; Stanton 2013); our interest here is in differentiating between terrorism and “small” forms of violence. Combining a conjoint experiment with machine learning techniques, our main findings are threefold.

First, the likelihood that ordinary citizens classify an event as terrorism is heavily dependent on relatively objective facts on the ground, such as the extremity and severity of violence employed. And yet, the public is also heavily influenced by the descriptions offered of who carried out the incident and why: acts are more likely to be seen as terrorism if they are carried out by organizations, less likely if they are carried out by individuals with histories of mental illness, more likely if they are carried out by Muslims, more likely if they are carried out in order to achieve political goals, and so on. Second, due to the subjective nature of some of these considerations, the media has significant agency in shaping how the public comes to classify violent events; given the extent to which terrorism increases news ratings, market-based theories of the media (e.g., Baum and Groeling 2010) would expect the media to have an incentive to frame ambiguous violent events in certain ways. Since violent incidents do not speak for themselves, our findings thus mirror other critiques of strictly “event-driven” theories of public opinion in foreign policy (e.g., Berinsky 2007). Third, although the public generally classifies events similarly to the formal classification schemes employed by different government agencies, it also deviates in some ways: violent incidents do not need to target civilians in order to be understood as terrorism, for example, and the public thinks that incidents motivated by hatred are just as likely to be terrorism as those motivated by more formal policy goals. In this sense, there seems to be a disconnect between formal legal definitions and our folk intuitions (Audi 2009) worth exploring.¹⁴

Future work should build on this analysis in four ways. First, although we show the powerful role that alternate frames can play in how and whether people classify events as terrorism, one of the interesting dynamics from a political science perspective is the extent to which

people are often presented with competing frames simultaneously (Druckman 2004): the *New York Post* covered the San Bernardino attacks very differently than MSNBC. Understanding how people classify events as terrorism in the face of competing frames is thus worthy of study. Second, although an automated content analysis of media coverage finds results consistent with our experimental findings, our theory also raises further questions about strategic media behavior: for example, if the media has ratings-based incentives to frame events consistent with the terrorist tropes discussed above, why emphasize white perpetrators’ histories of mental illness (Powell 2011), even though that lowers the likelihood of perceiving an event as terrorism? Additional analysis of media coverage will thus further enrich our understanding of how people come to understand events as terrorism.

Third, our results also suggest politicians can potentially manipulate perceptions of terrorism by framing violent incidents in certain ways. For example, if decision makers are in favor of pursuing more assertive foreign policies to combat a particular terrorist organization, we might expect them to highlight the potential for “foreign ties” in order to increase the likelihood the public perceives an incident to be terrorism and demand retribution. Finally, although our interest here is in how Americans think about terrorism, a final area of future research should explore whether foreign publics espouse similar judgments. The rise of knife attacks in Israel in 2015, for example, suggests a different portfolio of tactics would resonate with the public there; citizens of European countries with strict gun control laws might similarly be more likely to view shootings as terrorism than Americans do. We view each of these areas as an exciting part of future research that will continue to contribute to a deeper understanding of the important relationship between political violence and public opinion.

References

- Abrahms, Max. 2013. “The Credibility Paradox: Violence as a Double-Edged Sword in International Politics.” *International Studies Quarterly* 57(4): 660–71.
- Agamben, Giorgio. 2005. *State of Exception*. Chicago: University of Chicago Press.
- Albertson, Bethany, and Shana Kushner Gadarian. 2015. *Anxious Politics: Democratic Citizenship in a Threatening World*. New York: Cambridge University Press.
- Asal, Victor, and Aaron M. Hoffman. 2015. “Media Effects: Do Terrorist Organizations Launch Foreign Attacks in Response to Levels of Press Freedom or Press Attention?” *Conflict Management and Peace Science* 33(4): 381–99.

¹⁴As we show in Appendix §3 in the SI, a similar disconnect exists with scholarly definitions as well.

- Audi, Robert. 2009. *The Good in the Right: A Theory of Intuition and Intrinsic Value*. Princeton, NJ: Princeton University Press.
- Bali, Valentina A. 2007. "Terror and Elections: Lessons from Spain." *Electoral Studies* 26(3): 669–87.
- Baum, Matthew A., and Tim J. Groeling. 2010. *War Stories: The Causes and Consequences of Public Views of War*. Princeton, NJ: Princeton University Press.
- Berinsky, Adam J. 2007. "Assuming the Costs of War: Events, Elites, and American Public Support for Military Conflict." *Journal of Politics* 69(4): 975–97.
- Berrebi, Claude, and Esteban F. Klor. 2008. "Are Voters Sensitive to Terrorism? Direct Evidence from the Israeli Electorate." *American Political Science Review* 102(3): 279–301.
- Bloom, Mia. 2005. *Dying to Kill: The Allure of Suicide Terror*. New York: Columbia University Press.
- Braithwaite, Alex. 2013. "The Logic of Public Fear in Terrorism and Counter-terrorism." *Journal of Police and Criminal Psychology* 28(2): 95–101.
- Brewer, Marilynn B. 1999. "The Psychology of Prejudice: In-group Love or Outgroup Hate?" *Journal of Social Issues* 55(3): 429–44.
- Bueno de Mesquita, Ethan, and Eric S. Dickson. 2007. "The Propaganda of the Deed: Terrorism, Counterterrorism, and Mobilization." *American Journal of Political Science* 51(2): 364–81.
- Coady, Cecil A.J. 2004. "Terrorism and Innocence." *Journal of Ethics* 8(1): 37–58.
- Cowan, Richard, and Susan Cornwell. 2016. "Democrats Link Guns to Terrorism, Turn to Gun Control After Orlando." *Reuters*. <http://www.reuters.com/article/us-florida-shooting-guns-idUSKCN0Y22DV>
- Cox, Daniel and Robert P. Jones. 2015. "Nearly Half of Americans Worried That They or Their Family Will Be a Victim of Terrorism." *PRRI*. <https://www.prri.org/research/survey-nearly-half-of-americans-worried-that-they-or-their-family-will-be-a-victim-of-terrorism/>
- Crenshaw, Martha. 2000. "The Psychology of Terrorism: An Agenda for the 21st Century." *Political Psychology* 21(2): 405–20.
- D'Orazio, Vito, and Idean Salehyan. 2017. "Who Is a Terrorist? Ethnicity, Group Affiliation, and Understandings of Political Violence." Paper presented at the annual meeting of the International Studies Association.
- Druckman, James N. 2004. "Political Preference Formation: Competition, Deliberation, and the (Ir)relevance of Framing Effects." *American Political Science Review* 98(4): 671–86.
- Egami, Naoki, and Kosuke Imai. 2015. "Causal Interaction in High-Dimension." Working paper. <http://imai.princeton.edu/research/files/int.pdf>.
- Enders, Walter, and Todd Sandler. 2011. *The Political Economy of Terrorism*. New York: Cambridge University Press.
- Eubank, William Lee, and Leonard Weinberg. 1994. "Does Democracy Encourage Terrorism?" *Terrorism and Political Violence* 6(4): 417–35.
- Findley, Michael G., and Joseph K. Young. 2012. "Terrorism and Civil War: A Spatial and Temporal Approach to a Conceptual Problem." *Perspectives on Politics* 10(2): 285–305.
- Fortna, Virginia Page. 2015. "Do Terrorists Win? Rebels' Use of Terrorism and Civil War Outcomes." *International Organization* 69(3): 519–56.
- Getmansky, Anna, and Thomas Zeitzoff. 2014. "Terrorism and Voting: The Effect of Rocket Threat on Voting in Israeli Elections." *American Political Science Review* 108(3): 588–604.
- Goodwin, Jeff. 2006. "A Theory of Categorical Terrorism." *Social Forces* 84(4): 207–46.
- Hainmueller, Jens. 2012. "Entropy Balancing for Causal Effects: A Multivariate Reweighting Method to Produce Balanced Samples in Observational Studies." *Political Analysis* 20(1): 25–46.
- Hainmueller, Jens, Daniel J. Hopkins, and Teppei Yamamoto. 2014. "Causal Inference in Conjoint Analysis: Understanding Multidimensional Choices via Stated Preference Experiments." *Political Analysis* 22(1): 1–30.
- Hattem, Julian. 2015. "Clinton Calls Charleston Shooting 'Racist Terrorism.'" *The Hill*. <http://thehill.com/policy/national-security/245943-clinton-calls-charleston-shooting-racist-terrorism>
- Herrmann, Richard. 1988. "The Empirical Challenge of the Cognitive Revolution: A Strategy for Drawing Inferences about Perceptions." *International Studies Quarterly* 32(2): 175–203.
- Hodges, Adam. 2011. *The "War on Terror" Narrative: Discourse and Intertextuality in the Construction and Contestation of Sociopolitical Reality*. Oxford: Oxford University Press.
- Hoffman, Bruce. 2006. *Inside Terrorism*. New York: Columbia University Press.
- Horgan, John. 2004. *The Psychology of Terrorism*. New York: Routledge.
- Horowitz, Michael C. 2015. "The Rise and Spread of Suicide Bombing." *Annual Review of Political Science* 18: 69–84.
- Huddy, Leonie, Stanley Feldman, Charles Taber, and Gallya Lahav. 2005. "Threat, Anxiety, and Support of Antiterrorism Policies." *American Journal of Political Science* 49(3): 593–608.
- Huff, Connor, and Dominika Kruszewska. 2016. "Banners, Barricades, and Bombs: The Tactical Choices of Social Movements and Public Opinion." *Comparative Political Studies* 49(13): 1774–1808.
- Huff, Connor, and Dustin Tingley. 2015. "Who Are These People? Evaluating the Demographic Characteristics and Political Preferences of MTurk Survey Respondents." *Research & Politics* 2(3): 1–12.
- Imai, Kosuke, and Marc Ratkovic. 2013. "Estimating Treatment Effect Heterogeneity in Randomized Program Evaluation." *Annals of Applied Statistics* 7(1): 443–70.
- Kalyvas, Stathis N., and Laia Balcells. 2010. "International System and Technologies of Rebellion: How the End of the Cold War Shaped Internal Conflict." *American Political Science Review* 104(3): 415–29.
- Keller, Simon. 2005. *Philosophy 9/11: Thinking about the War on Terrorism*. Chicago: Open Court.
- Kertzer, Joshua D. 2007. "Seriousness, Grand Strategy, and Paradigm Shifts in the 'War on Terror.'" *International Journal* 62(4): 961–79.

- Kertzer, Joshua D., and Kathleen M. McGraw. 2012. "Folk Realism: Testing the Microfoundations of Realism in Ordinary Citizens." *International Studies Quarterly* 56(2): 245–58.
- Kertzer, Joshua D., Kathleen Powers, Brian C. Rathbun, and Ravi Iyer. 2014. "Do Moral Values Shape Foreign Policy Preferences?" *Journal of Politics* 76(3): 825–40.
- Kertzer, Joshua D., Jonathan Renshon, and Keren Yarhi-Milo. 2016. "How Do Observers Assess Resolve?" Working paper. http://people.fas.harvard.edu/~jkertzer/Research_files/KertzerRenshonYarhi-Milo021415.pdf
- Kydd, Andrew H., and Barbara F. Walter. 2006. "The Strategies of Terrorism." *International Security* 31(1): 49–80.
- Landler, Mark. 2013. "Obama Calls Blasts an 'Act of Terrorism.'" *New York Times*. <http://www.nytimes.com/2013/04/17/us/politics/obama-calls-marathon-bombings-an-act-of-terrorism.html>
- Lemieux, Anthony, and Elizabeth Karampelas. 2016. "‘Heil Hitler’ versus ‘Allahu Akbar’: An Experimental Approach to How Terrorism Is Differentially Perceived and Labeled." Paper presented at the annual meeting of the International Studies Association.
- Liberman, Nira, Yaacov Trope, and Elena Stephan. 2007. "Psychological Distance." In *Social Psychology: Handbook of Basic Principles*, eds. Arie W. Kruglanski and E. Tory Higgins. New York, NY: The Guilford Press, 353–83.
- Malle, Bertram F., and Joshua Knobe. 1997. "The Folk Concept of Intentionality." *Journal of Experimental Social Psychology* 33(2): 101–21.
- Merolla, Jennifer L., and Elizabeth J. Zechmeister. 2009. *Democracy at Risk: How Terrorist Threats Affect the Public*. Chicago: University of Chicago Press.
- Morgenthau, Hans J. 1985. *Politics among Nations: The Struggle for Power and Peace*. Brief ed. Boston: McGraw-Hill.
- Nacos, Brigitte L., Yaeli Bloch-Elkon, and Robert Y. Shapiro. 2011. *Selling Fear: Counterterrorism, the Media, and Public Opinion*. Chicago: University of Chicago Press.
- O'Rourke, Lindsey. 2009. "What's Special about Female Suicide Terrorism?" *Security Studies* 18(4): 681–718.
- Pape, Robert A. 2003. "The Strategic Logic of Suicide Terrorism." *American Political Science Review* 97(3): 343–61.
- Phillips, Brian J. 2014. "Terrorist Group Cooperation and Longevity." *International Studies Quarterly* 58(2): 336–47.
- Phillips, Brian J. 2015. "Was What Happened in Charleston Terrorism?" *The Monkey Cage*. <https://www.washingtonpost.com/news/monkey-cage/wp/2015/06/18/was-what-happened-in-charleston-terrorism/>.
- Powell, Kimberly A. 2011. "Framing Islam: An Analysis of U.S. Media Coverage of Terrorism since 9/11." *Communication Studies* 62(1): 90–112.
- Rodin, David. 2004. "Terrorism without Intention." *Ethics* 114(4): 752–71.
- Shanahan, Timothy. 2010. "Betraying a Certain Corruption of Mind: How (and How Not) to Define 'Terrorism.'" *Critical Studies on Terrorism* 3(2): 173–90.
- Smith, Megan, and James Igoe Walsh. 2013. "Do Drone Strikes Degrade Al Qaeda? Evidence from Propaganda Output." *Terrorism and Political Violence* 25(2): 311–27.
- Spaaij, Ramón. 2010. "The Enigma of Lone Wolf Terrorism: An Assessment." *Studies in Conflict & Terrorism* 33(9): 854–70.
- Stanton, Jessica A. 2013. "Terrorism in the Context of Civil War." *Journal of Politics* 75(4): 1009–22.
- Tilly, Charles. 1986. *The Contentious French*. Cambridge, MA: Harvard University Press.
- Uygur, Cenk. 2016. "Let's be Clear: If Muslims had Seized a Federal Building, they'd all be Dead by now. #whiteprivilege #OregonUnderAttack." *Tweet*. <https://twitter.com/cenkuygur/status/683770365009514497?lang=en>
- Weinberg, Leonard, Ami Pedahzur, and Sivan Hirsch-Hoefler. 2004. "The Challenges of Conceptualizing Terrorism." *Terrorism and Political Violence* 16(4): 777–94.
- Young, Joseph K. 2012. "What's in a Name? How to Classify Recent Violent Events." *Political Violence at a Glance*. <https://politicalviolenceataglance.org/2012/08/17/three-reasons-why-the-recent-attack-in-aurora-co-was-not-terrorism/>.
- Young, Joseph K., and Michael G. Findley. 2011. "Promise and Pitfalls of Terrorism Research." *International Studies Review* 13(3): 411–31.

Supporting Information

Additional Supporting Information may be found in the online version of this article at the publisher's website:

- Survey information
- Higher-order quantities of interest
- Scholarly classifications of terrorism
- Model diagnostics
- Alternative predicted probabilities
- Additional information about the media content analysis