

# HOW THE PUBLIC DEFINES TERRORISM

Connor Huff<sup>\*</sup>, Joshua D. Kertzer<sup>†</sup>

Last revised: July 4, 2016

**ABSTRACT:** Every time a major violent act takes place in the US, a public debate erupts as to whether it should be considered terrorism. Political scientists have offered a variety of conceptual frameworks, but have neglected to explore how ordinary citizens understand terrorism, despite the central role the public plays in our understanding of the relationship between terrorism and government action in the wake of violence. We synthesize components of both scholarly definitions and public debates to formulate predictions for how various attributes of incidents affect the likelihood they are perceived as terrorism. Using a conjoint experimental design, we show the importance of the extremity and severity of violence, but also the attributed motivation for the incident, and social categorization of the actor. The findings demonstrate how the language used to describe violent incidents, for which the media has considerable latitude, affects the likelihood the public classifies incidents as terrorism.

9846 words

---

<sup>\*</sup>Department of Government, Harvard University, Email: [cdezzanihuff@fas.harvard.edu](mailto:cdezzanihuff@fas.harvard.edu). Web: <http://connordhuff.com/>

<sup>†</sup>Department of Government, Harvard University, Email: [jkertzer@gov.harvard.edu](mailto:jkertzer@gov.harvard.edu). Web: <http://people.fas.harvard.edu/~jkertzer/>. We thank Bob Bates, Mia Bloom, Alex Braithwaite, Dara Kay Cohen, Naoki Egami, Marcus Holmes, Jeff Kaplow, In Song Kim, Melia Pfannenstiel, Peter Vining, Ariel White, Thomas Zeitzoff and audiences at the Political Violence Workshop at Harvard University, ISA 2016, MPSA 2016, and William and Mary for helpful comments, suggestions, and assistance. The authors acknowledge the support of the Weatherhead Center for International Affairs and Institute for Quantitative Social Science; thanks to Perry Abdulkadir, Michael Chen, and Aaron Miller for fantastic research assistance. Authors' names are in alphabetical order.

Following the mass shooting in Charleston, South Carolina, there was widespread debate throughout the United States about whether to classify the violent incident as terrorism. FBI Director James Comey determined that the shooting in Charleston was not terrorism, stating that “Terrorism is [an] act of violence...to try to influence a public body or citizenry, so it’s more of a political act. And again based on what I know so far I don’t see it as a political act.” His assessment was condemned from across the political spectrum, from critics on both the left (including then Democratic presidential candidate Hillary Clinton, and former Attorney General Eric Holder), and right (including then GOP presidential candidate Rick Santorum).<sup>3</sup> Similar public debates have erupted followed other violent incidents, including the bombing at the Boston Marathon and shooting in Orlando, Florida. These debates not only highlighted the contentiousness of classifying terrorism, but also the stakes involved in doing so, for policymakers, academics, and members of the public alike.

In this paper, we turn to experimental methods to explain the tenor of these public debates. We investigate terrorism in a public opinion context not because we believe that the mass public’s intuitions can necessarily resolve normative debates about what should or should not be considered terrorism, but rather because of the central role that public opinion plays in our understanding of how terrorism works. In a vast array of prior research, terrorism is understood as a form of violence that functions by attracting public attention. It is because terrorism hinges on public reaction that Margaret Thatcher suggested terrorists depend on “the oxygen of publicity,” that Carlo Pisacane declared terrorism to be “propaganda by deed,” and that Ayman al-Zawahiri suggested that for al-Qaeda, media coverage is “more than half” the battle (Smith and Walsh, 2013, 312). If the responses of ordinary citizens thus constitute a central causal mechanism through which terrorism operates, it logically follows that understanding *what ordinary citizens think terrorism is* is a crucial prerequisite to understanding *how* they react to it.

Employing a conjoint experiment on 1400 American adults, we show how ordinary citizens classify terrorism based not only on relatively objective facts on the ground, but also on fairly subjective considerations about the perpetrator. On the one hand, considerations about the extremity

---

<sup>3</sup> Clinton described the Charleston shooting as “an act of racist terror perpetrated in a house of God,” while Holder stated that Charleston was “clearly an act of terrorism,” and Santorum suggested that he didn’t “think there’s any question that this is an act of terrorism.” See <http://www.msnbc.com/msnbc/clinton-breaks-fbi-director-charleston-shooting>, [http://www.huffingtonpost.com/2015/07/09/james-comey-charleston-terrorism-\\_n\\_7764614.html](http://www.huffingtonpost.com/2015/07/09/james-comey-charleston-terrorism-_n_7764614.html), <http://www.thedailybeast.com/articles/2015/06/22/white-house-won-t-back-fbi-chief-on-charleston-terror.html>.

and severity of violence matter — though interestingly, not the distinction between civilian and military targets that plays a central role in contemporary legal definitions. On the other hand, our respondents are also heavily influenced by descriptions about the perpetrator’s identity and motivations, considerations about which the media has considerable latitude in the language it uses to cover incidents. In this sense, because actions do not always speak louder than words, the media has considerable power based on how it chooses to frame incidents: using predictive models derived from our experimental results, we show how the likelihood Americans will classify incidents with the characteristics of the attacks in San Bernadino, California, for example, can vary from around 30% to 82%, depending on the narrative constructed about the perpetrators’ identity and motivations (were they acting alone, or were their foreign ties real? Were they motivated by political goals, or personal ones? Are they described as having a history of mental illness?). In doing so we empirically demonstrate the ways in which terrorism can be socially constructed.

The discussion that follows has four parts. First, we discuss the important role that public opinion has played in our understanding of terrorism works, and present a simple typology of factors that either loom large in formal terrorism classification schemes, or in popular debates. Second, we discuss the design of our experiment. Third, we present our main results from our study, whose implications we highlight using machine learning techniques to construct counterfactual simulations of real world incidents, and automated content analyses of media coverage. We conclude by discussing the implications of our findings for the study of terrorism more generally.

## **The stakes of defining terrorism**

Whenever a major violent action takes place in the United States, a public debate erupts as to whether it should be classified as terrorism or not, mirroring a similarly contentious debate in the study of political violence as to how terrorism should be defined in the first place.

The contours of these debates matter for three sets of reasons. First, classifying actions as terrorism has direct policy implications for how the alleged perpetrators of the crime are prosecuted: in the United States, for example, terrorism is a federal charge, prosecuted in federal courts, unlike most other violent crime that is instead handled by the states. Similarly, transnational organizations

deemed to be terrorist entities are subjected to harsh financial sanctions, foreign nationals on terrorist watch lists are not allowed to board flights or enter the country, and Americans are prohibited from dealing with them. The consequences of acts being classified as terrorism are real and immediate. More generally, defining terrorism is crucial to fighting it: the United Nations has been unable to move forward with a comprehensive treaty against terrorism precisely because of the inability of its member states to come to a mutually acceptable definition on what terrorism is in the first place.

Second are important normative implications. As a discursive category, terrorism is understood as qualitatively different from other types of acts of violence: an extranormal or exceptional act, mandating an exceptional response (Agamben, 2005; Schmitt, 2008). In the 2004 election campaign, when John Kerry pronounced that terrorism was ultimately a law-enforcement issue, the George W. Bush campaign lampooned him in television advertisements for failing to take the threat seriously (Kertzer, 2007, pp. 964-5). Public debates following the shooting in Orlando in June 2016 displayed the same tendency. To categorize something as terrorism is to delegitimize its goals; terrorism isn't merely a problem to be *managed*, but one to be destroyed; terrorists are to be *hunted*, rather than negotiated with (Hodges, 2011). These normative questions are further exacerbated by complaints about the double standards with which the terrorist label is applied. When a group of armed white ranchers seized the headquarters of a US Fish and Wildlife Service complex in Oregon in January 2016 to protest the federal government's policies on grazing and land rights, left-leaning commentators condemned the refusal of the media to label the group as terrorists, and bemoaned the seeming leniency the group was being given by law enforcement: "if Muslims had seized a federal building, they'd all be dead by now", the commentator Cenk Uygur wrote, while hashtags like #YallQaeda, #YokelHaram and #VanillaISIS circulated on Twitter.<sup>4</sup> Scholars have similarly noticed a double standard, in which the media is more likely frame an attack as being linked to Islamic terror cells when the perpetrator is Muslim, and more likely to emphasize the attacker's personal life and mental instability if the perpetrator is not (Powell, 2011).

Third, these debates also have important stakes for political scientists. Almost every book or review article on terrorism we are aware of includes a compulsory line acknowledging the contentiousness of its definitional politics, and there are a variety of scholarly typologies and classi-

---

<sup>4</sup><http://www.rawstory.com/2016/01/yall-qaeda-twitter-users-mock-oregon-right-wing-militia-action-and-its-awesome/>.

fication schemes that political scientists regularly employ.<sup>5</sup> Indeed, a number of political scientists have used these typologies to constructively wade into public debates.<sup>6</sup> Just as our understanding of the dynamics of civil wars change when we adopt different operational definitions of the construct, our rules of thumb for how we know terrorism when we see it has important implications for which actions get turned into rows in our datasets, or cases in our analyses, and thus, our fundamental understanding of how terrorism works.

In this paper we shift the focus from how governments, the media, and academics classify terrorism, to instead explore how members of the mass public understand the term, as part of a growing body of research throughout the social sciences using experimental methods to unpack our folk intuitions about political concepts.<sup>7</sup> We investigate terrorism in a public opinion context not because we believe that our folk intuitions can resolve normative debates about what should and should not be considered terrorism — although disconnects between the judgments of citizens and the laws that govern them is indeed worthy of study (Audi, 2009) — but rather because of the central role that public opinion plays in our understanding of how terrorism works. Most of our academic models of terrorism emphasize “victim-target differentiation,” or the notion that terrorism is a communicative act (“propaganda by deed,” in the words of the Italian revolutionary Carlo Pisacane) directed not at the immediate victims of the attack, but broader audiences (Braithwaite, 2013; Fortna, 2015, p. 522; Asal and Hoffman, 2015). In Bueno de Mesquita and Dickson’s (2007) model of terrorism, terrorists carry out attacks in order to mobilize public support amongst the population they claim to represent; in Enders and Sandler (2011), terrorists seek to change government policy through public pressure, and in Kydd and Walter’s (2006) model of terrorism as a form of costly

---

<sup>5</sup>Indeed, more than twenty years ago Schmid and Jongman (1984, 1-38) noted that there were more than 100 different definitions of terrorism, while more recent work has noted that there are now over 200 definitions within the broader terrorism literature (Jackson, 2010). For examples of work that explicitly notes the contentiousness of defining terrorism, see Crenshaw 2000; Pape 2003; Horgan et al. 2004; Hoffman 2006; Kydd and Walter 2006; Merolla and Zechmeister 2009; Enders and Sandler 2011; Phillips 2014; Fortna 2015; Horowitz 2015; Huff and Kruszewska Forthcoming.

<sup>6</sup>For example, following the shooting in Charleston, Brian Phillips wrote an article in the Monkey Cage blog on the Washington Post titled “Was what happened in Charleston terrorism?” (<https://www.washingtonpost.com/blogs/monkey-cage/wp/2015/06/18/was-what-happened-in-charleston-terrorism/>). Similarly, following the shootings in Aurora, Colorado, and Oak Creek, Wisconsin, Joseph Young wrote a blog post on Political Violence at a Glance called “What’s in a Name? How to Classify Recent Violent Events.” (<http://politicalviolenceataglance.org/2012/08/17/three-reasons-why-the-recent-attack-in-aurora-co-was-not-terrorism/>.) In each of these pieces, the authors used typologies common to political scientists to help clarify whether particular events should be classified as terrorism.

<sup>7</sup>See e.g. Malle and Knobe, 1997; Kertzer and McGraw, 2012; Kertzer, Renshon and Yarhi-Milo, 2016.

signaling, three out of the five terrorist strategies they describe are directed at persuading the public.<sup>8</sup> The prominent role of public opinion is one mechanism frequently offered as to why terrorist groups turn to tactics like suicide bombing that tend to receive more publicity (Bloom, 2005), and why democracies are more likely to be targeted by terrorist attacks (Eubank and Weinberg, 1994; Pape, 2003; Stanton, 2013).<sup>9</sup>

Unlike many foreign policy issues, terrorism is highly salient to the public, capturing the public imagination — and producing downstream political consequences — to an extent that more ubiquitous forms of violence do not.<sup>10</sup> It receives extensive media coverage (Nacos, Bloch-Elkon and Shapiro, 2011), and fuels powerful emotional responses that not only affects the public’s attitudes towards foreign policy issues (Huddy et al., 2005; Albertson and Gadarian, 2015), but also how citizens act towards each other (Merolla and Zechmeister, 2009). Politicians are aware of terrorism’s sway on the popular imagination, which is one reason why Democratic lawmakers have begun using terrorism frames to pressure their Republican counterparts into supporting gun control measures. Outside the United States, political scientists have found that terrorist attacks increase support for right-wing parties in the locale where the attacks took place (Berrebi and Klor, 2008), decrease perceptions the government should be willing to negotiate with the perpetrator (Huff and Kruszewska, Forthcoming), sway elections (Bali, 2007), and that the mere threat of terrorism is enough to change voting behavior (Getmansky and Zeitzoff, 2014). Given this wealth of prior research in which public reactions play an important role as a causal mechanism linking violent incidents to downstream political consequences, it is crucial for political scientists to understand what the public thinks it is reacting to.

---

<sup>8</sup>It’s telling that even the conventional wisdom that Kydd and Walter are arguing against — that the purpose of terrorism is “to produce fear in an enemy population” also centers around public opinion.

<sup>9</sup>But see Chenoweth 2013.

<sup>10</sup>For example, a December 2015 PRRI poll found that 84% of Republicans and 70% of Democrats saw terrorism as a “critical issue” facing the country, as opposed to merely “one among many important issues.” In contrast, the same poll found that although 74% of Democrats called mass shootings a critical issue, only 60% of Republicans did. <http://www.prri.org/research/survey-nearly-half-of-americans-worried-that-they-or-their-family-will-be-a-victim-of-terrorism/>.

## A Typology for Classifying Terrorism

In this section we present a typology for classifying terrorism, illustrated in Figure 1. In so doing, we integrate common components of formal terrorism classification schemes with the contentious elements of public debates that erupt in the wake of violent incidents. The typology is comprised of two broad components. The first consists of relatively objective facts on the ground: information about the type and severity of the violence employed, and the incident’s target and location. This information is often available immediately after the incident occurs. In contrast, the second consists of information pertaining not to “what?” questions, but “who?” and “why?” These questions concern the identity of the perpetrators and motivation attributed to them. This category thus consists of information that is often far more subjective, often not available until days or weeks after an incident occurs, and which gives the media more leeway in how they frame events. In the remainder of this section we detail each of the components of the typology and derive a number of hypotheses.

Figure 1: A typology for classifying terrorist incidents

### **Objective characteristics of the action (“What?”)**

Violence	Target (military vs civilian)
Extremity (tactics)	Location (domestic vs international)
Severity (casualties)	

### **Subjective considerations about the actor (“Who?”) and motivation (“Why?”)**

Political purposiveness of actor: (individuals vs collectives)

Social categorization of actor (political and religious beliefs)

Motivation attributed to actor (politicization of objectives)

## Objective Criteria: The Facts on the Ground

The first component of the typology focuses on the fundamental descriptive characteristics of the action. Scholarly typologies and definitions of terrorism consistently emphasize two types of factors when discussing terrorist incidents. The first is *violence*, mentions of which usually appears in the first several words in almost every definition of the word of terrorism.<sup>11</sup> In the discussion below, we suggest violent incidents can be understood along two key dimensions: the extremity of the tactic used, and the severity of the consequences as measured by casualties. The second common component is information about whom the violence was against, information about which can be conveyed in either the type of target or location of the incident. We discuss each of these criteria in turn.

### Type and Severity of Violence

We argue that the violence observed in incidents can be understood along two dimensions. First, we can think of different “technologies” (Kalyvas and Balcells, 2010), or “repertoires of contention” (Tilly, 1986), characterized by the particular tactics employed in an incident. For example, bombing is a different type of violence than shootings or hostage-taking. Following the Boston Marathon Bombing in 2013, President Barack Obama stated that “Any time bombs are used to target innocent civilians, it is an act of terror.” This statement is emblematic of the close mapping in both academic and policy circles between the type of violence used and the likelihood the incident will be classified as terrorism, buttressed by recent work in political science on the role that extreme tactics play in structuring perceptions of actors (Abrahms, 2013; Huff and Kruszezwska, Forthcoming). The empirical implication is that some types of violence are more extreme than others, and that the more extreme the tactic, the more likely the incident will be classified as terrorism.

In addition to characterizing violence by its type, it can also be characterized by its *severity*. Indeed, the severity of the damage or magnitude of the carnage imposed by an attack is often the central focus of the media following violent incidents, with headlines typically bringing the number of casualties to the fore.<sup>12</sup> This leads to a second empirical implication: the higher the number of

---

<sup>11</sup>For example, the FBI defines terrorism as “the unlawful use of force and *violence*...”, and the State Department defines it as “premeditated, politically motivated *violence*.”

<sup>12</sup>For example, following the attacks in Paris in November 2015, an NBC news headline proclaimed “1 Amer-



casualties, the more likely the incident will be classified as terrorism.

### **The Target and Location**

We can characterize the target of violent incidents in two ways. First is whether the target is a civilian or non-combatant (Coady, 1985; Primoratz, 1990; Coady, 2004; Rodin, 2004). Consistent with the principle of distinction between combatants and non-combatants in *jus in bello*, the logic is that attacks upon the government or apparatus of the state might be undesirable, but is nonetheless legitimate in a way that targeting civilians is not. The clear empirical implication is that attacks on targets closely associated with formal state institutions are more likely to be considered legitimate, and thus should be less likely to be considered terrorism.

Second is the broader location in which the target is situated. This distinction was salient following the terrorist attacks in Paris and Lebanon in the fall of 2015, when a number of pundits explicitly criticized Western publics for the selective empathy revealed by their willingness to change their Facebook profile pictures in solidarity with Paris while ignoring the attacks which had occurred in Lebanon the previous day. Consistent with research on ingroup favoritism (Brewer and Brown, 1998) and affective responses decreasing with social and spacial distance (Liberman, Trope and Stephan, 2007, 373), the empirical implication of these critiques is that Americans should be more likely to classify events as terrorism when the incident takes place in the United States, compared to in foreign countries, particularly foreign countries unlike the United States.

### **Subjective Criteria: The Categorization of the Actor and Motivation**

If the previous definitional components focus on the *act* carried out, another set of components emphasizes the *actor* itself, and motivation understood as driving the actor's behavior. This idea is most commonly embedded in formal definitions employing verbiage noting that actors must have engaged in the act for "political aims and motives" (Hoffman, 2006), "achieving a political goal" (Keller, 2005), or "the advancement of some political, ideological, social, economic, religious, or

---

ican Among 129 Killed in Paris Terror Attacks," (<http://www.nbcnewyork.com/news/national-international/Paris-Attacks-France-Vows-Merciless-Response-348506901.html>), while following the San Bernadino attacks two weeks later, the CNN headline read "At least 14 people killed in shooting in San Bernadino; suspect identified" (<http://www.cnn.com/2015/12/02/us/san-bernardino-shooting/>).

military agenda” (Shanahan, 2010).

Three points are worth noting here. First, determining motivations is difficult, since it requires overcoming “the problem of other minds” and accessing an interior attribute of other actors, who may either have strategic reasons to misrepresent, or who themselves may not be aware of their underlying reasons for action. It is for this reason that classic IR scholars like Morgenthau (1985, 5) famously discouraged political scientists from studying motivation altogether. At the same time, however, as political scientists (Herrmann, 1988; O’Mahoney, 2015) and psychologists (Heider, 1958; Malle and Knobe, 1997) note, we intuitively rely on assumptions about motivations both when explaining action and evaluating the actors who carried them out.<sup>13</sup> In other words, the difficulty of assessing motivations does not deter us from relying on them.

Second, as a result, we often infer motivations for action using information about the actors themselves; we tend to provide answers to questions about “why?” using answers to questions about “who?”. It is thus worthwhile to parse out both analytically and empirically which particular components of descriptions of actors lead individuals to classify the acts the individuals carry out as terrorism. Third, precisely because of the subjectivity inherent in deriving meaning from actions, the media has a considerable amount of leeway in the narrative it constructs about the actor: from the characteristics of the actor it raises, to the rationale implied for the incident. We discuss each of these considerations in turn.

## Type of the Actor

Following violent incidents, governments, the media, and the public search for who is responsible, whether individuals or collectives. The differences are embodied in the distinction between “lone wolves” and “terrorist organizations.” We argue that we should be more likely to expect incidents perpetrated by organizations to be classified as terrorism than those by individuals. Because we tend to think of terrorism as “a collective, organized activity” (Spaaij, 2010, 855), and expect individuals to infer that formal organizations act strategically in pursuit of some broader political or social agenda, it follows that incidents carried out by lone individuals are less likely to be seen as terrorism

---

<sup>13</sup>See, for example, work in IR on the security dilemma, which is driven entirely by inaccurate assessments of the motivations of others (Jervis, 1978; Glaser, 1997).

than incidents carried out by groups, particularly more formal organizations. This means that when a bombing or hostage taking is carried out by an organization, for example, we tend to be more likely to assume it was carried out in pursuit of a broader political agenda than if the same incident is carried out by a “lone wolf.”

By contrast, in the wake of violent incidents perpetrated by individuals there is far more ambiguity about whether the action was taken purposively in pursuit of a broader agenda. Consider, for example, the wide-range of violence undertaken by individuals in the United States. Possessing the knowledge that an action was undertaken by an individual provides less information about whether an incident was undertaken purposively for some broader political agenda than whether an incident was perpetrated by an organization. In this sense, knowing information about the type of actor provides information about why the event occurred.

We can use similar logic to consider incidents perpetrated by individuals with a history of mental illness. These types of incidents stir some of the most controversial and public debates on when and how events should be defined as terrorism. For example, following the shootings perpetrated by Nidal Hassan at Fort Hood, Dylan Roof in Charleston, and Omar Mateen in Orlando, pundits debated whether the fact that these individuals had a history of mental illness should affect how we judge the actions they took. Just as we expect incidents perpetrated by individuals to be less likely to be classified as terrorism than those carried out by more formal organizations, we can similarly expect incidents perpetrated by individuals with a history of mental illness to be less likely to be considered terrorism. Because attacks by individuals with a history of mental illness are less likely to be attributed to a purposive political agenda, we expect their actions to be less likely to be considered terrorism.

### **Actor Categorization and Motivation**

The broader motivations for incidents are often understood in two ways. First, is through the social or identity categories used to describe the type of actor. Because of the pervasiveness of social categorization processes ([Brewer and Brown, 1998](#)), these descriptive words both provide information about the actor, and enable observers to draw inferences about broader political agendas. Some have argued that what matters is not only that the actor is perceived to be fighting for some broader

political agenda, but also, the particular agenda for which the actor is mobilized, hence the flurry of debates about the presence of “double standards” in who we define as terrorists. For example, it might matter if the actor is described as Muslim, rather than Christian, or left-wing, rather than right-wing. Most broadly, we hypothesize that incidents perpetrated by an actor described as Muslim are more likely to be classified as terrorism than other types of descriptions. Second is through the explicit motivation attributed for the action. For example, was the actor motivated by hatred toward the target, a personal agenda, or to change government policy? Each of these types of motivations provides varying levels of information about the purposiveness of the action, and thus, the likelihood we should expect an incident will be perceived as terrorism.

## Method

In this paper we seek to directly assess the attributes of incidents ordinary citizens use to define incidents as terrorism. We do so using a conjoint experimental design ([Hainmueller, Hopkins and Yamamoto, 2014](#); [Kertzer, Renshon and Yarhi-Milo, 2016](#)), in which we present our participants with a series of incidents with randomly generated features, and ask whether they would classify each incident as terrorism or not. The use of a ratings-based conjoint design offers two main advantages. First, it allows us to hold fixed a range of attributes of incidents which could confound observational studies on the topic. Even if there was an abundance of high-quality public opinion data on how the public defines terrorism, it would be extremely difficult to disentangle the effects of our attributes of interest, particularly given the presence of both incident-specific, and time-varying characteristics of the political environment surrounding the event, as well as our interest in the different types of media frames that prior research has shown vary with the type of actor. Given these problems, experiments give researchers vital control. Second, the use of a conjoint experimental design allows us to study the effect of each attribute using a relatively large number of experimental treatments, in a manner that would not be possible with more traditional factorial experiments due to statistical power constraints. While prior experimental research has sought to explore the effects of one or two of the factors we discuss in this paper on other dependent variables, a conjoint design allows us to

explore how they operate in tandem.<sup>14</sup>

## The Survey Instrument

In the experiment we manipulate seven attributes of an incident. As noted above, each of the attributes of the incident chosen, and the treatments within them, were included to reflect long-standing debates within both academic and policy worlds about how we should define terrorism. Each of these attributes is manipulated within a text block where the structure is intended to mirror what respondents would observe in the first few sentences of a newspaper article describing a recent incident. After respondents read about the details of the incident, they were then asked whether they would classify it as terrorism. Respondents go through this process seven times in total. Each of these attributes are presented in Table 1 and discussed in the text below.

First, we manipulated the extremity of the *tactics*: whether the actor carried out a protest, a hostage taking, a shooting, or a bombing. Protests were chosen as a baseline category of non-violent political action from which we would be able to test the hypothesis about how the use of more violent tactics subsequently affect whether individuals perceive an incident to be terrorism. The tactics of shooting and hostage taking were chosen to represent what we generally think of as violent incidents that may or may not be associated with formal organizations or terrorists. Finally, the tactic of bombing represents one of the most commonly used and emblematic tactics associated with terrorism. Each of these tactics also has the benefit in that each of them are actions that could plausibly be taken by a range of actors with varying ideologies, targets, and motivations.

Second, we manipulated the severity of the *casualties*, which was set to vary from either none, one, two, or ten. We specified these casualty levels with three goals in mind. First, as a way of unpacking the marginal effect of individuals dying as the result of an incident, to capture the notion that there is something systematically different about incidents that result in fatalities than those that do not. Second, to explore the functional form of the effect: is a single casualty enough for an incident to be seen as terrorism, or does the probability increase with severity? Finally, the numbers of casualties varied here is calibrated to reflect the generally few number of individuals killed in these types of incidents. In the Global Terrorism Database (GTD), for example, the median bombing has

---

<sup>14</sup>See Appendix §1.2 for a further discussion of rating-based conjoint.

(A) <i>Tactic</i>	The ... (1) ... protest (2) ... hostage taking (3) ... shooting (4) ... bombing
(B) <i>Target</i>	... occurred at a ... (1) ... military facility (2) ... police station (3) ... school (4) ... Christian community center (5) ... Muslim community center (6) ... Jewish community center (7) ... church (8) ... mosque (9) ... synagogue
(C) <i>Location</i>	... in ... (1) ... the United States. (2) ... a foreign democracy. (3) ... a foreign democracy with a history of human rights violations. (4) ... a foreign dictatorship. (5) ... a foreign dictatorship with a history of human rights violations.
(D) <i>Casualties</i>	There ... (1) ... were no individuals (2) ... was one individual (3) ... were two individuals (4) ... were ten individuals ... killed in the [Tactic].
(E) <i>Actor Description</i>	The [Tactic] was carried out by ... (1) ... an (2) ... a Christian (3) ... a Muslim (4) ... a left-wing (5) ... a right-wing
(F) <i>Actor Type</i>	... (1) ... organization. (2) ... organization with ties to the United States. (3) ... organization with ties to a foreign government. (4) ... group. (5) ... individual. (6) ... individual with a history of mental illness.
(G) <i>Actor Motivation</i>	News reports suggest ... (1) ... that there was no clear motivation for the incident. (2) ... the incident was motivated by the goal of overthrowing the government. (3) ... the incident was motivated by the goal of changing government policy. (4) ... the incident was motivated by hatred towards the target. (5) ... the individual had been in an ongoing personal dispute with one of the targets.

Table 1: Conjoint Study Treatments: Treatment categories are denoted by letters (A-G), while the numbers indicate the possible treatment sequences within each category. The treatment categories are connected into sentences using “...” to mirror the structure of the vignettes presented to respondents.

no fatalities; a bombing that kills 10 people is in the 96th percentile; a hostage-taking that kills 10 people is in the 98th percentile. We thus sought to mirror the real-world distribution of casualty levels for incidents commonly classified as terrorism, while also ensuring sufficient variation to be able to reflect higher casualty incidents.

Third, we manipulated the *target* of the incident, which was either a military facility, a police station, a school, a Christian community center, a Muslim community center, a Jewish community center, a church, a mosque, or a synagogue. The use of a military facility allows us to have a target that clearly signaled that it was not associated with civilians, while police stations reflect targets that are directly affiliated with the government but generally not perceived to be engaged in combat activities. The remaining seven categories represent different types of civilian targets, with the use of a school capturing a type of a target that is neither affiliated directly with combat activities nor religion. Finally, we distinguish between religious places of worship and religiously-affiliated community centers, due to the heightened symbolism of attacking the former.

Fourth, we manipulated the *location* of the incident, such that the act took place in either the United States, a foreign democracy, a foreign democracy with a history of human rights violations, a foreign dictatorship, or a foreign dictatorship with a history of human rights violations. We varied the location to explicitly test how differences in the attributes of the target population and government affect whether incidents should be perceived to be terrorism. By manipulating the location of the incident we are able to explicitly test whether and how incidents more closely linked to western institutions affect whether respondents are more likely to classify an event to be terrorism.

Fifth, we manipulated the *actor* who carried out the incident, varying whether the perpetrator was an organization, an organization with ties to the United States, an organization with ties to a foreign government, a group, an individual, or an individual with a history of mental illness. Employing these different types of perpetrators allows us to test the ways in which the perceived “purposiveness” of the actor affects the likelihood an incident is perceived to be terrorism.

Sixth, we manipulated the social *categorization* of the perpetrator, either describing the actor as Christian, Muslim, left wing, right wing, or a control condition in which no categorization was given, which acts as a baseline condition from which to interpret the effects of the other descriptions.

We chose these categories as they loomed the largest in recent prominent debates about classifying incidents as terrorism.

Finally, we manipulated the *motivation* attributed to the perpetrator: government overthrow, policy change, hatred towards the target, a personal dispute with the target, or no clear motivation. By including a category where the motivation is unclear, we can mirror the general uncertainty and absence of public information immediately following a violent incident, approximating the actual dissemination of information as an incident unfolds. The government overthrow and policy change motivations allow us to capture some of the most prominent goals of terrorist organizations as studied by political scientists (Kydd and Walter, 2006), while attributing the incident to hatred toward the target allows us to capture some of the more contentious incidents subject to recent debates. Finally, the use of a personal dispute as a baseline category allows us to study the effects of political motivations versus apolitical ones.

To be clear, our claim is not that the seven sets of factors discussed above and summarized in Table 1 constitute a complete set of characteristics that could be relevant to debates about terrorism, but rather that they mirror a wide set of characteristics from either public debates or the academic literature.<sup>15</sup> Indeed, despite not being an exhaustive list, we note that a fully-crossed factorial featuring all permutations of the factors we manipulate here would contain 108000 combinations.<sup>16</sup> An example of how the vignette appeared to respondents is presented below.

The incident: a shooting.

The shooting occurred at a church in a foreign democracy with a history of human rights violations.

There were two individuals killed in the shooting.

The shooting was carried out by a Muslim individual with a history of mental illness.

News reports suggest the individual had been in an ongoing personal dispute with one of the targets.

<sup>15</sup>Factors worth studying in future research include, for example, gender (O'Rourke, 2009).

<sup>16</sup>Thankfully, as noted above, conjoint experiments leverage between- and within-subject variation across participants to achieve higher levels of statistical power. Moreover, following best practices, we included randomization constraints to prevent implausible or problematic combinations of incident attributes from skewing the results. See Appendix §1.2.



## Fielding the Experiment

The experiment was embedded in a survey fielded on 1400 adults recruited using Amazon’s Mechanical Turk (MTurk) in August 2015.<sup>17</sup> MTurk has surged in popularity in recent years, because of the extent to which MTurk samples are “often more representative of the general population and substantially less expensive to recruit” than other convenience samples often used in political science (Berinsky, Huber and Lenz, 2012, 266). One commonly voiced concern with MTurk users is that while they are relatively diverse, they are not representative of the American population as a whole (Huff and Tingley, 2015). To mitigate against the possibility of biased estimates of our treatment effects, for the main analyses presented below we employ entropy balancing to reweight the data to population parameters (Hainmueller, 2012; for a similar empirical application, see Kertzer et al., 2014), though in Appendix §1.3 we find the substantive results remain the same regardless of whether weights are used. Because we have 1400 participants who classify 7 incidents each, the analyses that follow are of 9800 different randomly-generated scenarios, analyzed with clustered bootstrapped standard errors at the participant-level.

## Results

In this section we present the experimental results.<sup>18</sup> We do so in three parts. First, we demonstrate the importance of the type of tactic and number of casualties. In contrast, the target and location of the incident have no significant effect. Second, we show that the perceived political purposiveness of the actor, inferred either indirectly through their social categorization or directly through their posited motivation, has a significant effect on the likelihood an incident is classified as terrorism. In general, the perception that an actor is acting purposively for some broader political agenda matters more than the content of the agenda itself. We go on to demonstrate that there are significant

---

<sup>17</sup>See Appendix §1.4 for recruitment information, and a discussion of demand effects.

<sup>18</sup>In Appendix §1-§3, we present a range of additional methodological information and additional analyses, including additional information about our sample, a number of theoretically motivated interaction effects, and model diagnostics.

differences between liberals and conservatives, with conservatives being slightly more likely to classify incidents perpetrated by actors described as Muslim as terrorism.

## The Extremity of Tactic and Severity of the Violence

The results generally demonstrate that the *type* of violence plays a central role in determining whether an incident is classified as terrorism. In contrast, the *severity* of the violence seems to be much less salient in determining the classification. The results demonstrating the importance of tactics for determining whether the public defines an event as terrorism is presented in Figure 2. The left-hand panel of Figure 2 shows the relative importance of the tactics chosen by the actor on whether an incident is classified to be terrorism. As is standard in conjoint experiments, our quantities of interest here are Average Marginal Component Effects (AMCEs), depicted in a probability scale on the x-axis, such that point estimates further to the right indicate a greater probability that an event will be classified as terrorism. Two patterns are evident. First, the use of violence overwhelmingly increases the likelihood an incident will be classified to be terrorism. This is consistent with the prominent role violence plays throughout scholarly definitions and classification schemes. Second, the extremity of the violence matters as well, as bombings are significantly more likely than either shootings or hostage taking to be classified as terrorism. This finding is consistent with the statement made by President Obama in the wake of the Boston Marathon Bombing that implied a uniqueness around the tactic of bombing: this tactic is overwhelming associated with terrorist organizations. The results presented here demonstrate that this perception is similarly evident in ordinary citizens.<sup>19</sup>


The right-hand panel of Figure 2 shows the effect of increasing casualties on whether an incident is considered to be terrorism. While higher casualty incidents are more likely to be considered to be terrorism the size of the effect is small relative to the increase associated with more extreme tactics. In supplementary analyses in Appendix §2.2, we reiterate the point by showing that a bombing with no fatalities is statistically indistinguishable from a shooting that kills ten people.<sup>20</sup>

---

<sup>19</sup>A variety of mechanisms could explain why politicians and pundits perceive bombing to be a “more extreme” tactic than other types of violent tactics: the scope of their destructive capacity, their symbolic power, the necessary level of planning, and so on. We view the unpacking of the mechanisms underpinning perceptions of different types of violence to be an important area of future research.

<sup>20</sup>This finding is even more striking when considering the vast scope of violent incidents covered by this relatively small number of casualties: according to the Global Terrorism Database, 96% of terrorist bombings have fewer than 10 casualties. In this sense, the combinations of tactics and casualties we present here cover a substantial proportion

Figure 2: The effect of tactics and casualties



The figures demonstrate the importance of tactics in determining the likelihood that an incident is defined to be terrorism. Point estimates depicted with 95% CIs derived from  $B = 1500$  clustered bootstraps.

## The Target and Location

Figure 3 demonstrates that there is no significant relationship between either the target or location of an incident and the likelihood it is classified as terrorism. In particular, the results for the location of the incident show that there is no significant difference in the likelihood of an incident being classified as terrorism depending on whether the incident occurred in the United States or a foreign country. Moreover, there are no statistically significant differences based on regime type or histories of human rights violations. There is therefore little evidence that respondents employ differential definitions in contexts where the attack might be “justified” by prior government repression.


We observe similarly null effects for the target of the incident. That is, there is no statistically significant difference between how respondents view incidents targeting military facilities and other government, civilian, or religious centers. This absence is striking given the frequency with which terrorism definitions often distinguish between civilian and military targets; ordinary citizens seem not to make this distinction.

## Attributes of the Actor: Type, Descriptions, and Motivation

Thus far, we have shown that both the extremity and severity of the violent tactics employed strongly shape the likelihood individuals will classify an event as terrorism, while the target and location of the attack do not. The above factors are relatively objective “facts on the ground,” information available relatively shortly after an incident has taken place. In contrast, we now turn to relatively subjective characteristics about the actor itself about which the media has much greater latitude when calibrating coverage.

In Figure 4 we present the results from our actor treatment, which reveal two main patterns. First, acts carried out by collectives are around 15% more likely to be understood as terrorism than acts carried out by individuals. The effect is particularly strong for organizations with ties to foreign governments, showcasing the extent to which terrorism is thought of as politically purposive. Second, incidents perpetrated by an individual with a history of mental illness are significantly less likely than all other actors to have the incident classified as terrorism. This finding is in line with the of the violent incidents that loom large in public debates about terrorism.

Figure 3: Location and Target of incident.



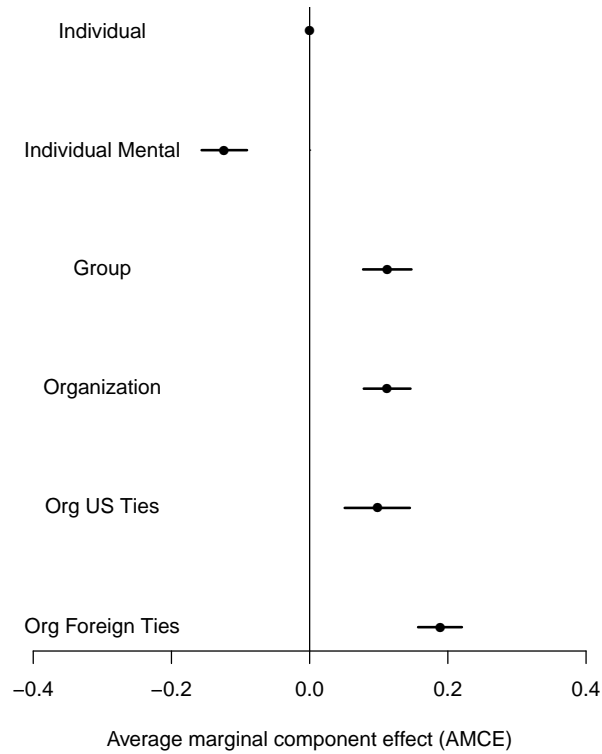
The AMCEs demonstrate that neither the location nor target are particularly important in determining the likelihood an incident will be defined to be terrorism.

pattern noted above: individuals with a history of a mental illness are less likely to be perceived to be acting purposively, and thus incidents they carry out are less likely to be understood as terrorism.

What is especially interesting about these findings is that some of these considerations are relatively subjective, thereby highlighting the way in which variation in the presentation of information about the actor affects whether it is perceived as terrorism. This means that the media has the ability to shape perceptions of violent incidents, and subsequently whether they are classified as terrorism, in the way that they choose to frame events: the decision to highlight the potential for links to either a formal organization or a foreign government (as evident in public debates over whether the San Bernadino attackers were truly affiliated with the Islamic State, or merely lone wolves “inspired” by them) or whether an individual has a history of mental illness, can significantly alter public perceptions of whether an incident is terrorism.

We see a less dramatic pattern in Figure 5, which depicts how the way in which the actor was described shapes perceptions of whether an incident should be classified as terrorism. The left-hand panel shows how providing a social categorization (particularly Muslim or Right-Wing) significantly increases the likelihood an incident will be classified as terrorism. The effect is strongest for perpetrators described as Muslim, though even then, the effect size remains relatively small compared to the other factors presented above. Interestingly, incidents where the perpetrator is described as Christian are not significantly more likely to be classified as terrorism. However, the right-hand panel shows that simply looking at the average effects of social categorizations belies important ideological gaps between liberals and conservatives. Two patterns are especially noteworthy. First, conservatives are more likely to classify incidents where the perpetrator is described as Muslim as terrorism than liberals. Second, liberals are more likely to classify incidents where the perpetrator is described as Christian as terrorism than conservative. Taken together, these findings highlight an interesting symmetrical interaction between respondents’ political ideology and the identity categories of the perpetrator. The fact that we find these effects in spite of potential social desirability biases suggests we should think of the magnitude of our Muslim effect as an under-, rather than over-estimate, although supplementary analyses reported in Appendix §2.3 generally reinforce that the effect of this bias is small.


Figure 4: The political purposiveness of the perpetrator.




The AMCEs demonstrate the importance of the actor being perceived as politically purposive. Incidents perpetrated by individuals with mental illness are significantly less likely to be perceived to be terrorism. Similarly, incidents perpetrated by groups, organizations, and organizations with ties to governments are more likely to be defined as terrorism than incidents perpetrated by individuals.

Figure 5: Social categorization effects

(a) Actor description




(b) Actor description for liberals and conservatives



The AMCEs demonstrate the the social categorization of the perpetrator is in general not particularly important in determining the likelihood an incident will be defined to be terrorism. However, there are notable differences between liberals (denoted by grey lines) and conservatives (denoted by black lines), the latter of whom are slightly more likely to classify incidents carried out by Muslims as terrorism.



Figure 6: Motive attribution effects



The AMCEs demonstrate that the more political the goals attributed to the organization, the more likely the incident is perceived to be terrorism.

Finally, Figure 6 presents the results for how the motivation attributed to the incident affects whether an incident is perceived to be terrorism. The main findings are threefold. First, incidents motivated by a personal dispute with the target are the least likely to be understood as terrorism. This matters given recent debates over whether, for example, the shooting of three Muslim university students in Chapel Hill, North Carolina in 2015 (ostensibly over a parking dispute) was terrorism or not. Second, compared to personal disputes, we find that incidents where the motivation is unclear increases the likelihood an incident is classified to be terrorism: when in doubt, our respondents are more likely to assume an incident is political than it is personal. Third, we find that from a baseline where the motivation is unclear, the motivations of hatred, policy change, and government overthrow increase the likelihood an incident is classified as terrorism. The powerful role of hatred here (an effect not statistically significant from that of policy change) is especially interesting since it does not exert a similarly large presence in legal definitions of terrorism, which tend to emphasize formal policy goals instead. Nonetheless, the strong effects of policy-oriented goals here reinforces the extent to which our participants associate terrorism with acting purposively for a broader political agenda.

## Extensions and applications

The strength of the conjoint experiment presented above is that it lets us examine the relative contributions of a large number of factors in shaping how ordinary citizens think about terrorism. However, it also has another virtue: we can invert our experimental findings to model how the public think about real world incidents. Doing so has three benefits. First, it allows us to code a range of prior incidents that have occurred in the real world using the typology presented above, to predict the probability that participants would classify each event as terrorism. Second, we exploit the subjectivity in many of these coding decisions to show how we can construct a number of different, equally plausible narratives about the same event, leading to sizable differences in the proportions of the public that should understand the event as terrorism or not. In this sense, we show empirically how terrorism can be socially constructed through media and elite framing. Finally, we can validate our findings using automated content analyses of media coverage. We sketch out all three, in turn.

## Application I: Mapping onto real world events

Table 2: The predicted probability a range of incidents are classified as terrorism using a weighted Support Vector Machine Classifier with Lasso Constraints.

Incident	Date	Predicted Probability
East Selma Church Shooting	(09/20/15)	0.23
Dallas Police HQ Shooting	(06/13/15)	0.25
UCLA Black Lives Matter Protest	(10/08/15)	0.28
University of California Tuition Hike Protests	(03/18/15)	0.31
Zvornik Police Station shooting	(04/27/15)	0.33
Marysville Pilchuck High School shooting	(10/24/14)	0.33
Rocori High School shooting	(09/24/03)	0.33
Shooting of Police in Brooklyn	(12/20/14)	0.35
Islamic Community Center of Phoenix Demonstrations	(10/10/15)	0.35
Shooting of Police in Oakland	(03/21/09)	0.35
Pentagon Metro Shooting	(03/04/10)	0.36
St. Columbanus Church Shooting	(11/26/12)	0.39
Poe Elementary School Bombing	(09/15/59)	0.39
Newport Church hostage situation	(07/30/06)	0.39
Copenhagen Synagogue Shooting	(02/14/15)	0.39
University of Alabama Huntsville	(12/02/10)	0.43
Charleston church Shooting	(06/17/15)	0.46
Camp Shelby Shootings	(08/05/15)	0.47
Overland Park Jewish Community Center Shooting	(04/13/14)	0.47
Knoxville Unitarian Universalist Church shooting	(07/27/08)	0.48
Rosemary Anderson High School shooting	(12/13/14)	0.49
Shooting of George Tiller	(05/31/09)	0.55
Bombing of Shiraa village mosque	(12/30/14)	0.55
Kehilat Bnei Torah synagogue attack	(11/18/14)	0.59
KKK Selma Bombing	(09/15/63)	0.61
Seattle Jewish Federation Shooting	(07/28/06)	0.62
Nag Hammadi massacre	(01/07/10)	0.62
Contra attack in Quilali	(11/11/87)	0.62
ETA Sanguesa car bombing	(05/30/03)	0.64
Lombard Islamic School bombing	(08/12/12)	0.71
Hamas attack on IDF in Khan Yunis	(12/24/14)	0.71
Shebaa Farms incident	(01/28/15)	0.72
Camp Integrity Suicide bombing	(08/07/15)	0.74
Porte de Vincennes hostage situation	(01/09/15)	0.74
Pakistan Army General HQ hostage situation	(10/10/09)	0.78
Zif School Bombing	(09/17/02)	0.84
Aksu bombing	(08/19/10)	0.85
Chattanooga shootings	(07/16/15)	0.85
Fort Hood Shootings	(11/05/09)	0.85

First, we code 39 recent or otherwise prominent incidents that can be subsumed by the typology presented above. Because of the potential non-additive effects of the attributes we explore here,

we employ a machine learning method proposed by [Imai and Ratkovic \(2013\)](#) and [Egami and Imai \(2015\)](#), who apply a variable selection problem approach to the study of high-dimensional treatment interactions by adapting a Support Vector Machine classifier with LASSO constraints, thereby estimating a high-dimensional interactive model without overfitting the data.<sup>21</sup> We then calculate fitted values to estimate the predicted probability that each incident would be understood as terrorism, presented in Table 2.

Several patterns in Table 2 are striking, including the extent to which bombings perpetrated by formal organizations (such as Hamas, Hezbollah, ETA, and the Ku Klux Klan) are the types of incidents most likely to be perceived as terrorism, as well as the proportion of incidents with predicted probabilities clustering around 50%, illustrating the contentiousness and difficulty of defining terrorism in the wake of violent incidents.

## Application II: Demonstrating framing effects

Especially important for our purposes, though, is that the ranking of many of these incidents is in part a function of relatively subjective coding decisions regarding the description of the actor and motivation. In our view, this is one of the most striking and important implications of this paper, since the media make similar coding decisions when choosing whether to highlight a suspect’s religion, speculate on foreign ties, attribute a political motivation behind the incident, or psychoanalyze the perpetrator from a distance.

To make this discussion concrete, consider the shooting which took place on December 2nd, 2015, in San Bernadino, California. There are at least three subjective decisions one can make about how to code the event. First, was the act carried out by lone individuals, or were they actors with foreign ties? Much of the media coverage after the attack concerned Tashfeen Malik, the female shooter, pledging allegiance to ISIS on Facebook during the attack. Is pledging allegiance on social media enough to constitute a foreign tie? Second, was the attack motivated by policy goals, or was the motivation behind the shooting unclear? Third, is shooting up a community center indicative of mental illness, or merely malign intentions? Our model suggests these coding decisions

---

<sup>21</sup>Consistent with the main results presented throughout this paper we present weighted estimates here; see Appendix §4 for alternative modeling specifications.

have important implications for how people understand the attack. If the perpetrators are coded as having foreign ties, and with the goal of changing policy, our model suggests an 81.6% likelihood it will be understood as terrorism. If, instead, the perpetrators are instead simply seen as individuals with no clear motivation, the probability of classifying the incident as terrorism drops to 50.3%. If the relevant narrative instead raises the specter of mental illness, and doesn't mention the attackers' religion, the probability drops to 30.6%. In this sense, there is over a 50% difference in the likelihood the attacks will be classified as terrorism, based solely on how the story is framed.

### Application III: Analyzing Media Coverage

The results presented in Table 2 also suggest two additional implications. First, if our classifications are relatively accurate, they should correlate with the rate at which the incidents are described as terrorism in the real world: an event with the characteristics that our participants see as highly likely to be terrorism should generally be described as such in the press, for example. Second, the extremity of the probability estimates should be negatively correlated with the magnitude of public debate: events clustered at around 50% on our probability scale should be the ones that produce the greatest dissensus.


Automated content analyses of newspaper coverage thus serve as a helpful means of validating our findings. As a plausibility probe, we selected two events from Table 2: an incident with a high classification rate (the shooting at Fort Hood, with a terrorism probability of 85%), and an incident with a classification rate close to 50% (the shooting at Charleston, with a terrorism probability of 46%), and used LexisNexis to collect every newspaper article about the incidents published in the two weeks following each attack.<sup>22</sup> There are several important similarities between these incidents: for example, both shootings received high levels of media coverage, were perpetrated by an individual, and took place in the United States. Yet there are also a number of important differences which map directly onto the typology presented in this paper: for example, the perpetrator in Charleston had a more pronounced history of mental illness, was associated with right-wing extremists, and was described as being motivated by hatred rather than formal policy change.

---


<sup>22</sup>While we ideally would have been able to also compare these incidents with another with a low predicted probability, the low predicted-probability events generally did not garner enough media attention; we thus focus our analysis on two shootings that received a sufficient level of media coverage to be able to conduct meaningful analyses.

Figure 7: Automated content analysis of media coverage

(a) Proportion of articles mentioning “terrorism” or “terrorist”



(b) Proportion of articles mentioning “debate”



Automated content analysis of daily newspaper coverage of the attacks in Charleston (which our model classifies as being 86% likely to be classified as terrorism) and Fort Hood (which our model classifies as 46% likely to be classified as terrorism) provides results consistent with our theory. Panel a shows that newspaper coverage of the Fort Hood attack is consistently more likely to refer to terrorism than articles covering the Charleston attack. Panel b shows that newspaper coverage of Charleston was significantly more likely to discuss the presence of “debate” than newspaper coverage of Fort Hood.

For each event we calculated two quantities of interest. The first is the proportion of articles published making reference to “terrorism” or “terrorist”: if our model is accurate, we should expect that news coverage of the incident with a high predicted classification rate should more frequently invoke terrorism than the coverage of the event with a moderate predicted classification rate. The second is the proportion of articles published making reference to “debate”: if our model is reliable, we should expect that the coverage of the event with a high predicted classification rate should feature less debate than the coverage of the event with a predicted classification rate that hovers near a coin toss. Sure enough, when we plot the daily counts of both quantities in Figure 7, we find that articles about the shooting at Fort Hood were more likely to mention the words “terrorism” or “terrorist” than articles about the shooting in Charleston. Indeed, in the two-week window around each of the respective incidents, articles about Charleston were approximately 13% less likely to mention the word “terrorism” or “terrorist” than articles about Fort Hood.<sup>23</sup> This finding is consistent with the idea that incidents we predict as more likely to be classified as terrorism are indeed discussed as such in the news media. Similarly, the right panel of Figure 7 demonstrates that articles about the shooting at Charleston were more likely to mention the word “debate” than articles about the shooting in Fort Hood. In the two-week window around the incidents, articles about Charleston were 15% more likely to mention the word “debate” than articles about Fort Hood.<sup>24</sup> Although more of a plausibility probe than a full test of our theory, these results show how media coverage offers another domain in which our theoretical framework can be explored.

## Conclusion

This paper contributes to broader research on the relationship between public opinion and government action in the wake of violent events by exploring the characteristics that shape how ordinary citizens classify incidents as terrorism. A rich literature in IR explores the distinction between terrorism and “big” forms of violence like civil wars (e.g. Findley and Young, 2012; Stanton, 2013); our interest here is differentiating between terrorism and “small” forms of violence. Combining a

---

<sup>23</sup>The percentage of articles mentioning “terrorism” or “terrorist” for Fort Hood is 31% while for Charleston it was 18%. The difference in proportion is statistically significant at the  $\alpha = 0.01$  level.

<sup>24</sup>The percentage of articles mentioning “debate” for Fort Hood is 5% while for Charleston it was 20%.

conjoint experiment with machine learning techniques, our main findings are threefold.

First, the likelihood that ordinary citizens classify an event as terrorism is heavily dependent on relatively objective facts on the ground such as the extremity and severity of violence employed. And yet, the public is also heavily influenced by the descriptions offered of who carried out the incident and why: acts are more likely to be seen as terrorism if they are carried out by organizations, less likely if they are carried out by individuals with histories of mental illness, more likely if they are carried out by Muslims, more likely if they are carried out in order to achieve political goals, and so on. Second, due to the subjective nature of some of these considerations, this means that the media has significant agency in shaping how the public comes to classify violent events; given the extent to which terrorism increases news ratings, market-based theories of the media (e.g. [Baum and Groeling 2010](#)) would expect the media to have an incentive to frame ambiguous violent events in certain ways. Since violent incidents do not speak for themselves, our findings thus mirror other critiques of strictly “event-driven” theories of public opinion in foreign policy (e.g. [Berinsky, 2007](#)). Third, although the public generally classifies events similarly to the formal classification schemes employed by different government agencies, it also deviates in some ways: violent incidents do not need to target civilians in order to be understood as terrorism, for example, and the public thinks that incidents motivated by hatred are just as likely to be terrorism as those motivated by more formal policy goals. In this sense, there seems to be a disconnect between formal legal definitions and our folk intuitions ([Audi, 2009](#)) worth exploring.

Future work should build on this analysis in four ways. First, although we show the powerful role that alternate frames can play in how and whether people classify events as terrorism or not, one of the interesting dynamics about these incidents from a political science perspective is the extent to which people are often presented with competing frames simultaneously ([Druckman, 2004](#)): the *New York Post* covered the San Bernadino attacks very differently than MSNBC. Understanding how people classify events as terrorism in the face of competing frames is thus worthy of study. Second, although an automated content analysis of media coverage finds results consistent with our experimental findings, our theory also raises further questions about strategic media behavior: for example, if the media has ratings-based incentives to frame events consistent with the terrorist tropes



discussed above, why emphasize white perpetrators' histories of mental illness (Powell, 2011), even though that lowers the likelihood of perceiving an event as terrorism? Carrying out additional analysis of media coverage will thus further enrich our understanding of how people come to understand events as terrorism.

Third, our results also point to the ways in which politicians can potentially manipulate perceptions of terrorism by framing violent incidents in certain ways. For example, if decision-makers are in favor of pursuing more aggressive and activist foreign policies to combat a particular terrorist organization, we might expect them to highlight the potential for "foreign ties" in order to increase the likelihood the public perceives an incident to be terrorism and demand retribution. Finally, while our interest here is on how Americans think about terrorism, a final area of future research would be to explore whether foreign publics espouse similar judgments. The rise of knife attacks in Israel, for example, suggests a different portfolio of tactics would resonate with the public there; citizens of European countries with strict gun control laws might similarly be more likely to view shootings as terrorism than Americans do. We view each of these four areas as an exciting area of future research that will continue to contribute to a deeper understanding of the important relationship between political violence and public opinion.

## References

- Abrahms, Max. 2013. "The Credibility Paradox: Violence as a Double-Edged Sword in International Politics." *International Studies Quarterly* 57(4):660–671.
- Agamben, Giorgio. 2005. *State of Exception*. Chicago, IL: University of Chicago Press.
- Albertson, Bethany and Shana Kushner Gadarian. 2015. *Anxious Politics: Democratic Citizenship in a Threatening World*. New York, NY: Cambridge University Press.
- Asal, Victor and Aaron M. Hoffman. 2015. "Media effects: Do terrorist organizations launch foreign attacks in response to levels of press freedom or press attention?" *Conflict Management and Peace Science* Forthcoming.
- Audi, Robert. 2009. *The good in the right: A theory of intuition and intrinsic value*. Princeton University Press.
- Bali, Valentina A. 2007. "Terror and elections: Lessons from Spain." *Electoral Studies* 26(3):669–687.
- Baum, Matthew A. and Tim J. Groeling. 2010. *War Stories: The Causes and Consequences of Public Views of War*. Princeton, NJ: Princeton University Press.

- Berinsky, Adam J. 2007. "Assuming the Costs of War: Events, Elites, and American Public Support for Military Conflict." *Journal of Politics* 69(4):975–97.
- Berinsky, Adam J, Gregory A Huber and Gabriel S Lenz. 2012. "Evaluating online labor markets for experimental research: Amazon. com's Mechanical Turk." *Political Analysis* 20(3):351–368.
- Berrebi, Claude and Esteban F Klor. 2008. "Are voters sensitive to terrorism? Direct evidence from the Israeli electorate." *American Political Science Review* 102(03):279–301.
- Bloom, Mia. 2005. *Dying to kill: The allure of suicide terror*. Columbia University Press.
- Braithwaite, Alex. 2013. "The Logic of Public Fear in Terrorism and Counter-terrorism." *Journal of Police and Criminal Psychology* 28(2):95–101.
- Brewer, Marilynn B. and Rupert J. Brown. 1998. Intergroup Relations. In *The Handbook of Social Psychology*, ed. Daniel T. Gilbert, Susan T. Fiske and Gardner Lindzey. New York: McGraw-Hill pp. 554–594.
- Bueno de Mesquita, Ethan and Eric S Dickson. 2007. "The propaganda of the deed: Terrorism, counterterrorism, and mobilization." *American Journal of Political Science* 51(2):364–381.
- Chandler, Jesse, Pam Mueller and Gabriele Paolacci. 2014. "Nonnaïveté among Amazon Mechanical Turk workers: Consequences and solutions for behavioral researchers." *Behavior research methods* 46(1):112–130.
- Charnysh, Volha, Christopher Lucas and Prerna Singh. 2014. "The Ties That Bind National Identity Salience and Pro-Social Behavior Toward the Ethnic Other." *Comparative Political Studies* p. 0010414014543103.
- Chenoweth, Erica. 2013. "Terrorism and democracy." *Annual Review of Political Science* 16:355–378.
- Coady, Cecil A.J. 1985. "The morality of terrorism." *Philosophy* 60(231):47–69.
- Coady, Cecil A.J. 2004. "Terrorism and innocence." *The Journal of Ethics* 8(1):37–58.
- Crenshaw, Martha. 2000. "The Psychology of Terrorism: An Agenda for the 21st Century." *Political Psychology* 21(2):405–420.
- Druckman, James N. 2004. "Political Preference Formation: Competition, Deliberation, and the (Ir)relevance of Framing Effects." *American Political Science Review* 98(4):671–686.
- Egami, Naoki and Kosuke Imai. 2015. "Causal Interaction in High-Dimension." Working paper.  
**URL:** <http://imai.princeton.edu/research/files/int.pdf>
- Enders, Walter and Todd Sandler. 2011. *The political economy of terrorism*. Cambridge University Press.
- Eubank, William Lee and Leonard Weinberg. 1994. "Does democracy encourage terrorism?" *Terrorism and Political Violence* 6(4):417–435.
- Findley, Michael G and Joseph K Young. 2012. "Terrorism and civil war: A spatial and temporal approach to a conceptual problem." *Perspectives on Politics* 10(02):285–305.
- Fortna, Virginia Page. 2015. "Do Terrorists Win? Rebels' Use of Terrorism and Civil War Outcomes." *International Organization* pp. 519–56.
- Frederick, Shane. 2005. "Cognitive reflection and decision making." *The Journal of Economic Perspectives* 19(4):25–42.

- Gaines, Brian J., James H. Kuklinski and Paul J. Quirk. 2007. "The Logic of the Survey Experiment Reexamined." *Political Analysis* 15(1):1–20.
- Getmansky, Anna and Thomas Zeitzoff. 2014. "Terrorism and voting: The effect of rocket threat on voting in Israeli elections." *American Political Science Review* 108(03):588–604.
- Glaser, Charles L. 1997. "The Security Dilemma Revisited." *World Politics* 50(1):171–201.
- Grice, H. Paul. 1975. Logic and Conversation. In *Syntax and Semantics*, ed. Peter Cole and Jerry L. Morgan. Vol. 3 New York, NY: Academic Press pp. 41–58.
- Hainmueller, Jens. 2012. "Entropy Balancing for Causal Effects: A Multivariate Reweighting Method to Produce Balanced Samples in Observational Studies." *Political Analysis* 20(1):25–46.
- Hainmueller, Jens and Daniel J Hopkins. 2015. "The hidden american immigration consensus: A conjoint analysis of attitudes toward immigrants." *American Journal of Political Science* 59(3):529–548.
- Hainmueller, Jens, Daniel J Hopkins and Teppei Yamamoto. 2014. "Causal inference in conjoint analysis: Understanding multidimensional choices via stated preference experiments." *Political Analysis* 22(1):1–30.
- Healy, Andrew and Gabriel S Lenz. 2014. "Substituting the End for the Whole: Why Voters Respond Primarily to the Election-Year Economy." *American Journal of Political Science* 58(1):31–47.
- Heider, Fritz. 1958. *The Psychology of Interpersonal Relations*. New York: Wiley.
- Herrmann, Richard. 1988. "The Empirical Challenge of the Cognitive Revolution: A Strategy for the Drawing Inferences about Perceptions." *International Studies Quarterly* 32(2):175–203.
- Hodges, Adam. 2011. *The "War on terror" narrative: discourse and intertextuality in the construction and contestation of sociopolitical reality*. Oxford University Press.
- Hoffman, Bruce. 2006. *Inside terrorism*. Columbia University Press.
- Horgan, John et al. 2004. *The psychology of terrorism*. Routledge.
- Horowitz, Michael C. 2015. "The Rise and Spread of Suicide Bombing." *Annual Review of Political Science* 18:69–84.
- Huddy, Leonie, Stanley Feldman, Charles Taber and Gallya Lahav. 2005. "Threat, anxiety, and support of antiterrorism policies." *American journal of political science* 49(3):593–608.
- Huff, Connor and Dominika Kruszezwska. Forthcoming. "Banners, Barricades, and Bombs: The Tactical Choices of Social Movements and Public Opinion." *Comparative Political Studies*.
- Huff, Connor and Dustin Tingley. 2015. "'Who are these people?' Evaluating the demographic characteristics and political preferences of MTurk survey respondents." *Research & Politics* 2(3).
- Imai, Kosuke and Marc Ratkovic. 2013. "Estimating Treatment Effect Heterogeneity in Randomized Program Evaluation." *The Annals of Applied Statistics* 7(1):443–470.
- Imhoff, Roland and Julia Recker. 2012. "Differentiating Islamophobia: Introducing a new scale to measure Islamoprejudice and secular Islam critique." *Political Psychology* 33(6):811–824.
- Jackson, Richard. 2010. "An argument for terrorism." *Perspectives on terrorism* 2(2).
- Jervis, Robert. 1978. "Cooperation Under the Security Dilemma." *World Politics* 30(2):167–214.

- Kalyvas, Stathis N. and Laia Balcels. 2010. "International System and Technologies of Rebellion: How the End of the Cold War Shaped Internal Conflict." *American Political Science Review* 104(3):415–429.
- Keller, Simon. 2005. *Philosophy 9/11: thinking about the war on terrorism*. Open Court Publishing.
- Kertzer, Joshua D. 2007. "Seriousness, grand strategy, and paradigm shifts in the" war on terror". *International Journal* pp. 961–979.
- Kertzer, Joshua D., Jonathan Renshon and Keren Yarhi-Milo. 2016. "How Do Observers Assess Resolve?" Working paper.  
**URL:** [http://jonathanrenshon.com/Site/CurrentResearch\\_files/KertzerRenshonYarhi-Milo021415.pdf](http://jonathanrenshon.com/Site/CurrentResearch_files/KertzerRenshonYarhi-Milo021415.pdf)
- Kertzer, Joshua D. and Kathleen M. McGraw. 2012. "Folk Realism: Testing the Microfoundations of Realism in Ordinary Citizens." *International Studies Quarterly* 56(2):245–258.
- Kertzer, Joshua D., Kathleen Powers, Brian C. Rathbun and Ravi Iyer. 2014. "Do moral values shape foreign policy preferences?" *Journal of Politics* 76(3):825–840.
- Kriner, Douglas L and Francis X Shen. 2013. "Reassessing American Casualty Sensitivity The Mediating Influence of Inequality." *Journal of Conflict Resolution* p. 0022002713492638.
- Krupnikov, Yanna and Adam Seth Levine. 2014. "Cross-sample comparisons and external validity." *Journal of Experimental Political Science* 1(01):59–80.
- Kydd, Andrew H and Barbara F Walter. 2006. "The strategies of terrorism." *International Security* 31(1):49–80.
- Liberman, Nira, Yaacov Trope and Elena Stephan. 2007. "Psychological distance." *Social psychology: Handbook of basic principles* 2:353–383.
- Malle, Bertram F. and Joshua Knobe. 1997. "The folk concept of intentionality." *Journal of Experimental Social Psychology* 33(2):101–121.
- Merolla, Jennifer L and Elizabeth J Zechmeister. 2009. *Democracy at risk: How terrorist threats affect the public*. University of Chicago Press.
- Morgenthau, Hans J. 1985. *Politics Among Nations: The Struggle for Power and Peace*. Brief ed. Boston, MA: McGraw-Hill.
- Mutz, Diana C. 2011. *Population-Based Survey Experiments*. Princeton University Press.
- Nacos, Brigitte L, Yaeli Bloch-Elkon and Robert Y Shapiro. 2011. *Selling fear: Counterterrorism, the media, and public opinion*. University of Chicago Press.
- O'Mahoney, Joseph. 2015. "Why did they do that?: the methodology of reasons for action." *International Theory* 7(2):231–262.
- O'Rourke, Lindsey. 2009. "What's Special about Female Suicide Terrorism?" *Security Studies* 18(4):681–718.
- Paolacci, Gabriele and Jesse Chandler. 2014. "Inside the turk understanding mechanical turk as a participant pool." *Current Directions in Psychological Science* 23(3):184–188.
- Pape, Robert A. 2003. "The strategic logic of suicide terrorism." *American Political Science Review* 97(03):343–361.

- Phillips, Brian J. 2014. "Terrorist group cooperation and longevity." *International Studies Quarterly* 58(2):336–347.
- Powell, Kimberly A. 2011. "Framing Islam: An analysis of US media coverage of terrorism since 9/11." *Communication Studies* 62(1):90–112.
- Primoratz, Igor. 1990. "What is terrorism?" *Journal of Applied Philosophy* 7(2):129–138.
- Rodin, David. 2004. "Terrorism without Intention." *Ethics* 114(4):752–771.
- Schmid, Alex P and Albert J Jongman. 1984. *Political Terrorism: A Research Guide to Concepts, Theories, Databases and Literature*. North Holland.
- Schmitt, Carl. 2008. *The concept of the political: Expanded edition*. University of Chicago Press.
- Shanahan, Timothy. 2010. "Betraying a certain corruption of mind: how (and how not) to define 'terrorism'." *Critical Studies on Terrorism* 3(2):173–190.
- Smith, Megan and James Igoe Walsh. 2013. "Do Drone Strikes Degrade Al Qaeda? Evidence From Propaganda Output." *Terrorism and Political Violence* 25(2):311–327.
- Spaaij, Ramón. 2010. "The enigma of lone wolf terrorism: An assessment." *Studies in Conflict & Terrorism* 33(9):854–870.
- Stanton, Jessica A. 2013. "Terrorism in the Context of Civil War." *The Journal of Politics* 75(04):1009–1022.
- Tilly, Charles. 1986. *The Contentious French*. Cambridge, MA: Harvard University Press.
- Tomz, Michael and Jessica Weeks. 2013. "Public opinion and the democratic peace." *American political science review* 107(4):849–865.
- Tversky, Amos and Daniel Kahneman. 1973. "Availability: A heuristic for judging frequency and probability." *Cognitive Psychology* 5(2):207–232.
- Wallace, Geoffrey PR. 2013. "International law and public attitudes toward torture: An experimental study." *International Organization* 67(01):105–140.

# HOW THE PUBLIC DEFINES TERRORISM

## *Supplementary Appendix*

### Contents

<b>1</b>	<b>Survey information</b>	<b>3</b>
1.1	Terrorism salience and the timing of the survey . . . . .	3
	Figure 1: Terrorism salience as measured by newspaper coverage and Google search data . . . . .	4
1.2	Conjoint randomization details . . . . .	5
1.3	Survey weighting . . . . .	6
	Table 1: Sample characteristics . . . . .	6
	Figure 2: Weighted and unweighted Average Marginal Component Effects (AMCEs). . . . .	8
1.4	MTurk protocol and demand effects . . . . .	9
<b>2</b>	<b>Higher-order quantities of interest</b>	<b>10</b>
2.1	Average Marginal Treatment Interaction Effects (AMTIEs) . . . . .	10
	Figure 3: Estimated Ranges of AMTIEs (I) . . . . .	11
	Figure 4: Estimated Ranges of AMTIEs (II) . . . . .	12
2.2	Interaction Effects of Interest . . . . .	13
	Figure 5: Interaction between the motivation and tactics. . . . .	14
	Figure 6: Interaction between the tactic and actor. . . . .	15
	Figure 7: Interaction between the actor and their motivation. . . . .	16
	Figure 8: Interaction between the tactic and the number of casualties. . . . .	18
	Figure 9: Interaction between the tactic and the target. . . . .	19
	Figure 10: Interaction between the type of actor and their ideology. . . . .	20
2.3	The moderating role of attitudes towards Islam . . . . .	21
	Figure 11: AMCEs by levels of Islamoprejudice . . . . .	23
	Figure 12: AMCEs by levels of Secular Critique against Islam. . . . .	24
<b>3</b>	<b>Model diagnostics</b>	<b>25</b>
3.1	Randomization checks . . . . .	25
	Table 2: Randomization check for tactic. . . . .	25
	Table 3: Randomization check for casualties. . . . .	25
	Table 4: Randomization check for description of the actor. . . . .	25
	Table 5: Randomization check for the actor ideology. . . . .	25
	Table 6: Randomization check for the motivation of the actor. . . . .	25
	Table 7: Randomization check for the target of the incident. . . . .	25
	Table 8: Randomization check for the location of the incident. . . . .	26
3.2	Carryover assumption test . . . . .	26
	Figure 13: AMCEs display little evidence of longitudinal effects . . . . .	27

<b>4</b>	<b>Alternative predicted probabilities</b>	<b>28</b>
	Table 9: Predicted probabilities: Support Vector Machine (SVM) classifier . . . . .	29
	Table 10: Predicted probabilities: Linear Probability Model (LPM) . . . . .	30
	Table 11: Predicted probabilities: Linear Probability Model (LPM) with weights . .	31
<b>5</b>	<b>Additional information about the media content analysis</b>	<b>32</b>

# 1 Survey information

## 1.1 Terrorism salience and the timing of the survey

Given the potential for pretreatment effects (Gaines, Kuklinski and Quirk, 2007) and availability bias (Tversky and Kahneman, 1973), it is worthwhile to contextualize the survey experiment’s timing. After all, a survey experiment fielded in the days after the 9/11 attacks would likely report different understandings of terrorism than one fielded in the 1950s. While our purpose is to show how ordinary citizens currently define terrorism rather than extrapolate to other time periods, it is nonetheless valuable to explore how close the survey fielding was to highly salient incidents that could be affecting the results.

We thus measured the salience of terrorism in two different ways. First, we looked at media coverage of terrorism in a twelve month window around our experiment, recording daily counts of every newspaper article on Lexis Nexis mentioning “terrorism” or “terrorist” from between February 17, 2015 and February 17, 2016.<sup>25</sup> Second, we tracked public interest in terrorism using data from Google Trends, recording the popularity of the term “terrorism” in Google searches in the United States from January 2015 to January 2016.

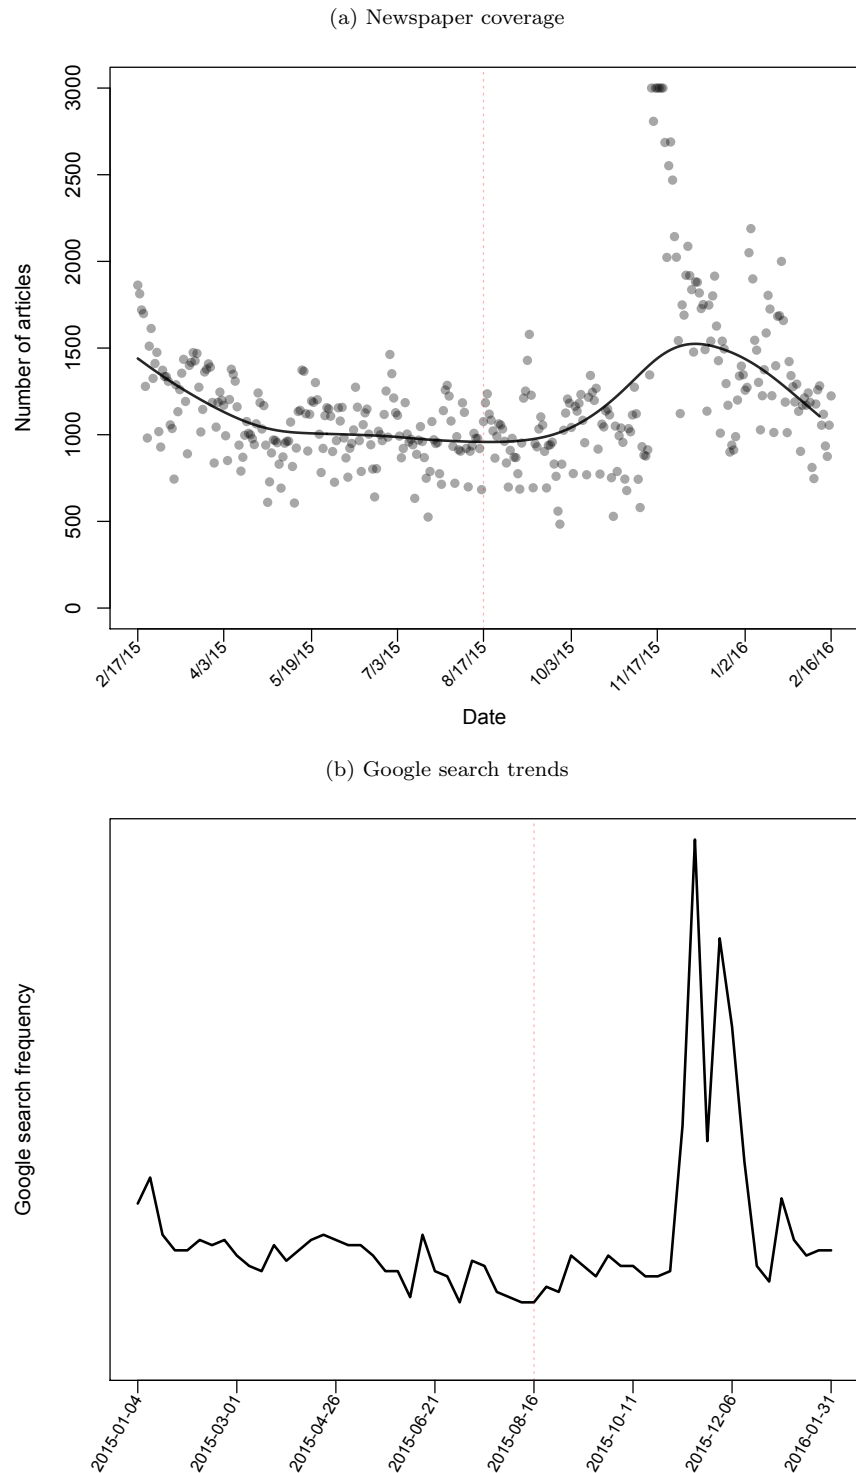
Figure 1(a) depicts the salience of terrorism as measured by media coverage, with each point on the scatterplot indicating the daily number of newspaper articles mentioning “terrorism” or “terrorist”, overlaid with a nonparametric smoothing spline with seven knots. Figure 1(b) depicts the salience of terrorism as measured by weekly Google search trends. In both plots, the dashed vertical line in the middle of the plot indicates the fielding of our study. As it turns out, both plots suggest roughly the same story: although there is a significant spike in salience corresponding with the Paris attacks in mid-November 2015, both newspaper coverage and Google search trends remain relatively constant throughout the rest of the year. Importantly, then, the plots show that the study’s fielding did not coincide with a highly salient event that could potentially bias the findings.

---

<sup>25</sup>To avoid overcounting, we removed article duplicates with high similarity. The search for these articles was conducted on Lexis Nexis on June 25, 2016. The distance of the search from the publication of the articles alleviates concerns that there might be differential indexing of articles which could plausibly vary across the search window.



Figure 1: Terrorism salience as measured by newspaper coverage and Google search data



The top panel depicts daily counts of newspaper articles mentioning terrorism within six months of our experiment's fielding; the dark black line depicts a natural smoothing spline with seven knots. The bottom panel shows Google trends data for terrorism from January 2015-January 2016. The dashed vertical line in each plot indicates the date of our study's fielding. Both plots show stark temporary increases in terrorism salience following the Paris attacks, but that otherwise salience is fairly constant (including during the period in which the study was fielded), mitigating concerns about bias in our results. Note that because of limits on the number of results that can be displayed by Lexis-Nexis for any given time period, the results are truncated at 3000 for six of the seven days following the Paris attacks.

## 1.2 Conjoint randomization details

Following best practices with conjoint experiments (e.g. [Hainmueller and Hopkins, 2015](#); [Kertzer, Renshon and Yarhi-Milo, 2016](#)), we constrained the randomization to prevent implausible or problematic combinations of incident attributes from skewing the results. First, we restricted scenarios where there were casualties from the protests. This was intended to avoid scenarios where it was unclear whether the government or actor were responsible for perpetrating the casualties, thereby avoiding conflating terrorism perpetrated by non-state actors from that perpetrated by governments.<sup>26</sup> Second, we prevented organizations with ties to the United States from attacking the United States. Third, we prevented organizations with ties to the United States from having a left-wing, Muslim, or Christian ideology. Fourth, we restricted all actors that were not individuals (organizations, groups, etc.) from being motivated by an ongoing personal dispute with one of the targets. This was done to capture the purposeful nature with which organizations act in their targeting. If, for example, an individual that is a member of an organization engaged in violent activity associated with a personal dispute this action would generally not be attributed to their organization. Finally, we precluded the combination of protests being motivated by a personal dispute with the target.

We should note that our experimental protocol differs from traditional conjoint experiments in two ways. First, following [Hainmueller and Hopkins \(2015\)](#), conjoint experiments in political science have thus far tended to be *choice-based* conjoints, in which participants are presented between pairs of profiles, and asked to choose between them. However, because of the purpose of our study — for which it would be problematic to present participants with two incidents, and have them decide which one is “more terrorist” than the other — we instead employed a *rating-based* conjoint, in which participants are presented with only one profile at a time, and simply asked whether they consider the incident to be terrorism or not. In this way, the the experimental design more closely approximates the phenomenon in the world we are attempting to study. Second, conjoint experiments that use tabular formats randomize the order in which attributes are presented at the respondent-level to test for order effects. Because we use an article format instead, our pretesting found that randomizing the order in which attributes were presented suffered from problems of lexical interpretability: narratives have structure, such that being told why something took place before what the thing itself was, for example, made the scenarios hard to follow. We therefore fixed the order in which the attributes were presented to the one that was the most interpretable for our participants. Although this precludes the possibility of testing for order effects, it improves our inferences by more closely resembling the

---

<sup>26</sup>This is not to say that we view heinous actions taken by the government to be more or less likely to be terrorism — only that it raises a number of sufficiently complex methodological issues that we view it as an exciting area of future research.

linguistic conventions our participants are familiar with (Grice, 1975).

### 1.3 Survey weighting

One of the main goals of this paper is to make inferences about how ordinary members of the public classify events as terrorism. Since our participants come from a convenience sample recruited through Amazon Mechanical Turk, it is important to think carefully about how respondents on MTurk might be fundamentally different from other types of individuals in the broader population. Of course, professional polling firms like YouGov and SSI often face a similar challenge: they have a finite population of respondents who have agreed to take their surveys, and hope to make inferences to a broader target population such as the U.S. public. A common technique used to approximate the U.S. public involves weighting respondents on known population parameters, up-weighting responses from individuals in demographic categories underrepresented in the sample.<sup>27</sup> This weighting technique is commonly used in prominent survey research in political science, such as the Cooperative Congressional Election Survey (CCES).<sup>28</sup>

Table 1: Sample characteristics

Sample Characteristic	Unweighted Sample	Weighted Sample	Population Target
Female	0.449	0.459	0.508
Male	0.527	0.518	0.492
18 to 24 years	0.124	0.140	0.128
25 to 44 years	0.659	0.375	0.342
45 to 64 years	0.191	0.351	0.341
65 years and over	0.025	0.134	0.189
High School or less	0.129	0.376	0.420
Some college	0.282	0.212	0.194
College/University	0.487	0.309	0.282
Graduate/Professional school	0.102	0.103	0.104

In this paper we use a similar weighting technique, using entropy balancing (Hainmueller, 2012) to reweight our sample towards standard population parameters, as presented in Table 1, trimming the weights to reduce the impact of extreme values.<sup>29</sup> Because the addition of weights only change quantities of interest in survey experiments if the demographic characteristics used to generate the weights significantly interact with the study treatments, adding weights to survey experiments


<sup>27</sup>Additionally, because of our interest in whether liberals and conservatives define terrorism differently, we employed a two-stage sampling strategy in which we fielded an initial version of the survey, and then produced a second version that was only visible to MTurk workers who had self-identified as conservative in an earlier, unrelated study fielded by one of the authors, thereby partially mitigating the liberal nature of the MTurk worker pool (Huff and Tingley, 2015). This allowed us to ensure that we had a sufficient number of respondents in each of the demographic categories upon which we would subsequently weight.

<sup>28</sup>For a discussion of the weights used by professional polling firms and how the types of individuals on these platforms are similar to those on MTurk, see Huff and Tingley (2015).

<sup>29</sup>See Kertzer et al. (2014) for another application of this procedure.

generally does not substantially change the results ([Mutz, 2011](#)), and as shown in Figure 2, our results here are no exception. The weighted estimates are substantively similar to those presented in the paper, and the results do not substantively differ regardless of whether weights are used. The only treatment category for which the results slightly differ from those in the main text is for the location of the incident. However, as is discussed in the main body of the text, the substantive size of the effect for the location is extremely small relative to the other effect sizes, such that the overall findings remain the same.

Figure 2: Weighted and unweighted Average Marginal Component Effects (AMCEs).



Grey lines indicate weighted estimates while black lines are unweighted.

## 1.4 MTurk protocol and demand effects

As noted in the main text, participants in the study were recruited on Amazon Mechanical Turk (MTurk). Participants were paid \$1.50 for their participation. We also limited participation in the study to MTurk workers located in the United States, who had completed greater than 50 HITs, and whose HIT approval rate was greater than 95%. MTurk is increasingly popular for experimental research in political science, and experimental studies conducted on MTurk have been published in a variety of notable journals, including the *American Political Science Review* (Tomz and Weeks, 2013), *American Journal of Political Science* (Healy and Lenz, 2014), *Comparative Political Studies* (Charnysh, Lucas and Singh, 2014), *International Organization* (Wallace, 2013), and the *Journal of Conflict Resolution* (Kriner and Shen, 2013). Nonetheless, there are two potential concerns about the use of MTurk. The first involves the composition of the sample, which we address through the use of survey weights and entropy balancing, detailed above. Additionally, because of our interest in whether liberals and conservatives define terrorism differently, we employed a two-stage sampling strategy in which we fielded an initial version of the survey, and then produced a second version that was only visible to MTurk workers who had self-identified as conservative in an earlier, unrelated study fielded by one of the authors, thereby partially mitigating the liberal nature of the MTurk worker pool (Huff and Tingley, 2015).

The second concerns demand effects, in that like participants in other online survey platforms, MTurk users often participate in a large number of studies, which is problematic for studies that either require naive participants (e.g. the cognitive reflection test - Frederick, 2005), or employ commonly used experimental paradigms (Chandler, Mueller and Paolacci, 2014; Paolacci and Chandler, 2014; Krupnikov and Levine, 2014). Given both the intended purpose of our study and relative novelty of our experimental design, however, this concern does not apply here.

## 2 Higher-order quantities of interest

### 2.1 Average Marginal Treatment Interaction Effects (AMTIEs)

The analysis in the main text presents Average Marginal Component Effects (AMCEs), the standard quantity of interest in conjoint experiments (Hainmueller, Hopkins and Yamamoto, 2014). Since the interpretation of interaction effects in conjoint experiments is sensitive to the choice of the baseline category, Egami and Imai (2015) propose an alternate quantity of interest, Average Marginal Treatment Interaction Effects (AMTIEs). Figures 3-4 thus present the full-range of one-way, two-way, and three-way AMTIEs for each of the treatment categories presented in the paper, letting us test for the possibility of higher-order interactive effects between sets of treatments.

Figure 3(a) replicates the conclusions presented in the main text: tactics display the substantively largest effects in determining whether ordinary citizens classify events as terrorism or not; substantively large effects are also detected for motivation, and the categorization of the actor; casualties also have substantively important effects, but at a lesser magnitude. In contrast, the ideology of the actor, and the target and location have relatively minor effects. Two points are worth emphasizing here. First, as in the main text, we see that subjective considerations about the perpetrator matter more than many objective facts on the ground. Second, since the one-way AMTIE is equivalent to an AMCE for each of the treatment categories as a whole (rather than the levels within them), we should not be surprised we find equivalent results with either quantity.

Figure 3(b) presents the complete set of two-way AMTIEs. Most of these two-way interactions have relatively small effects; only the two largest interactions have effect ranges greater than 0.1. However, four of the five largest two-way AMTIE ranges involve tactics, and two of these effects (the interaction between tactics and motivation, and between tactics and the categorization of the actor) are substantively large, which we explore in greater detail in the next section. Similarly, Figure 4(c) shows that the three-way effects are extremely small, although as before, the AMTIEs with the largest effect ranges tend to involve tactics (which make up the thirteen substantively largest three-way effects). Finally, Figure 4(d) summarizes the three previous panels by presenting the AMTIEs for the four largest effect ranges across the one-, two-, and three-way interaction categories. In the next section, we unpack these different interactive effects.

Figure 3: Estimated Ranges of AMTIEs (I)







Figure 4: Estimated Ranges of AMTIEs (II)



## 2.2 Interaction Effects of Interest

The previous set of analyses illustrated in Figures 3-4 suggest that although most of the treatments in the experiment lack substantively large interaction effects, there are a number of interaction effects worth exploring in greater detail. We thus begin by presenting interaction effects for the substantively largest two-way AMTIEs from Figure 3(b), before turning to a broader set of interaction effects of theoretical interest.

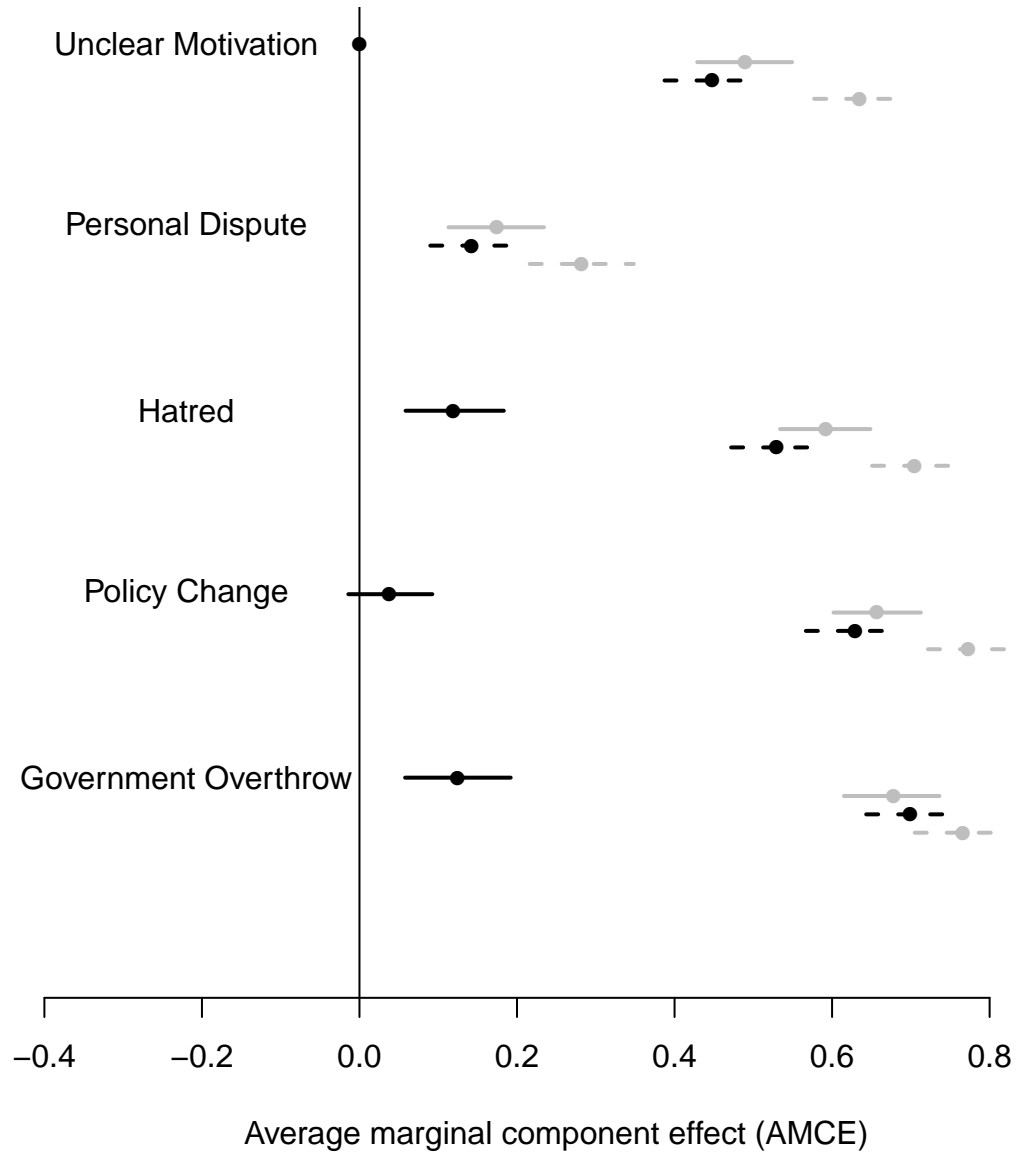
Figure 5 presents the results for the interaction between the motivation of the actor and the tactics employed in the incident, the substantively largest effect range of the two-way AMTIEs depicted in Figure 3(b). Several findings are worth noting. First, and most importantly, the results reinforce the extent to which violence is seen as a necessary factor for an event to be classified as terrorism: protests are unlikely to be classified as terrorism, regardless of the motivation attributed to the actor. Second, for incidents with an unclear motivation, moving from non-violence (protests) to a violent tactic increased the likelihood an incident is classified as terrorism by over 40%. Third, incidents perpetrated by actors motivated by personal disputes are infrequently considered to be terrorism. This is consistent with the understanding of terrorism as reflecting a broader political agenda.

Figure 6 presents the results for the interaction between the actor and the tactic employed in the incident, the second-largest effect range of the two-way AMTIEs depicted in Figure 3(b). The results demonstrate that much of the effect heterogeneity comes from a mental illness “discount”: violent incidents perpetrated by individuals with a history of mental illness are less likely to be classified as terrorism than incidents perpetrated by all other actors, regardless of the tactics employed. The results also demonstrate that while protests are similarly unlikely to be classified as terrorism, violent incidents perpetrated by more formal organizations are much more likely to be classified as terrorism than those carried out by individuals. This is indicative of the perception discussed in the paper that collectives of individuals are more likely to be acting in pursuit of a broader political agenda.

Figure 7 presents the results for the actor and the motivation of the actor, the two-way AMTIE from Figure 3(b) with the third-largest effect range. The results here reinforce the above discussion about the discounting effect of mental illness in classifying events as terrorism. Although more overtly political motivations generally lead to a greater likelihood of perceiving an incident as terrorism, the effect is considerably compressed for individuals described as having a history of mental illness.

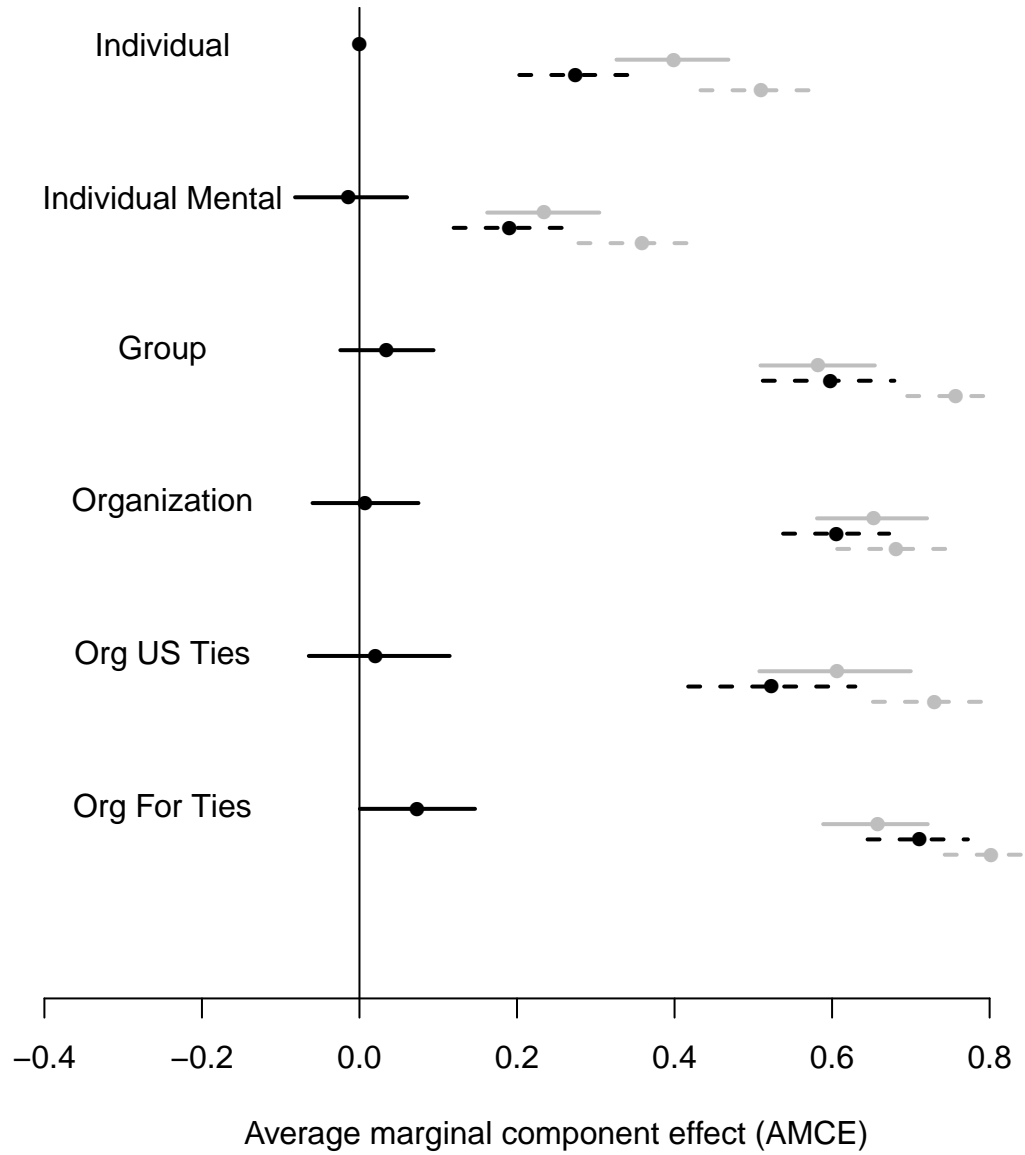
Figure 8 presents the results for the tactic and the number of casualties. The results here reiterate the point that while high casualty incidents are more likely to be considered terrorism the effect size is small when compared to the increase associated with more extreme tactics. For example, the results demonstrate that a bombing with zero casualties is more likely to be considered to be

Figure 5: Interaction between the tactics employed and motivation attributed to the actor.



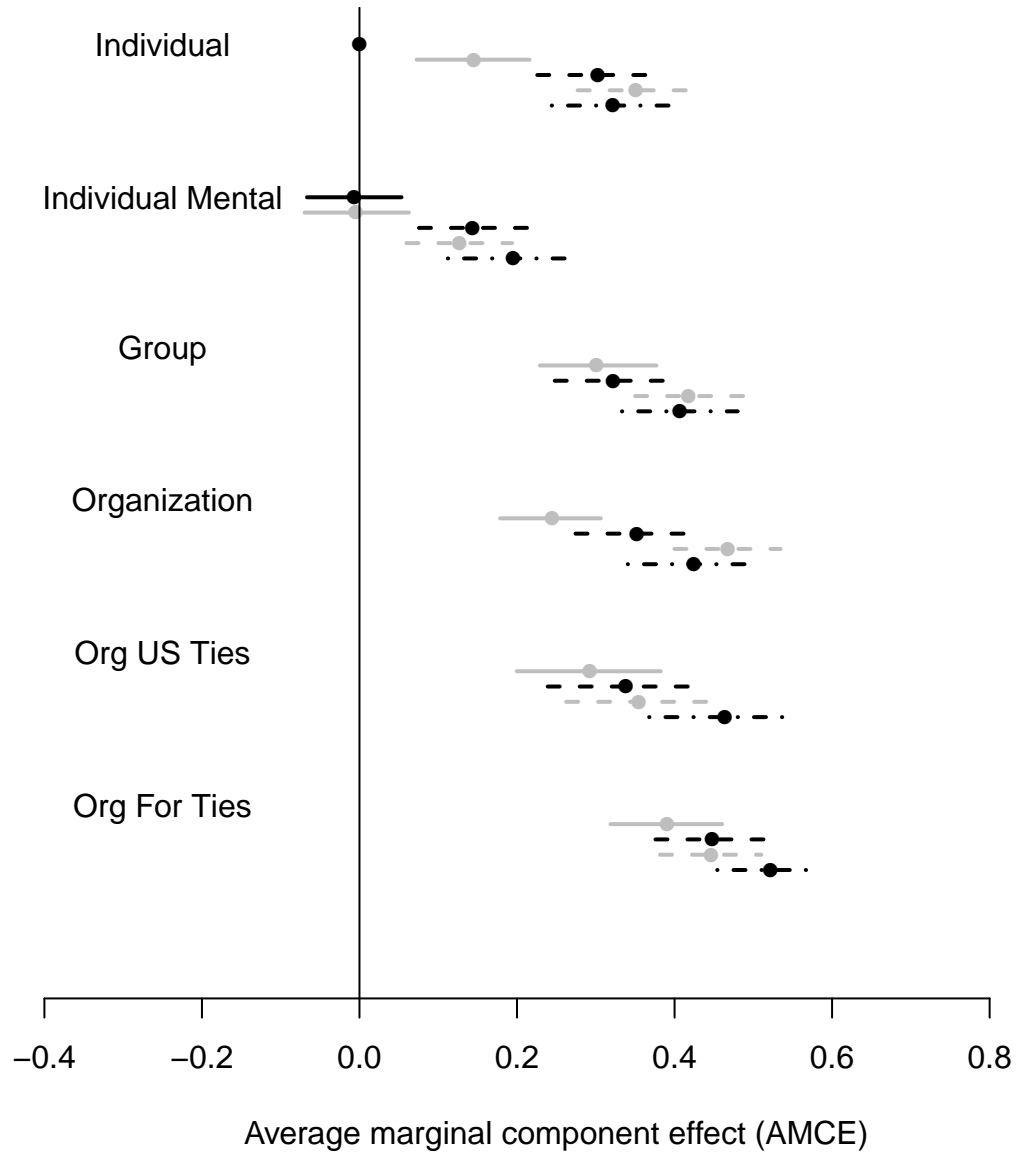
Interaction between the motivation and tactics. Black lines indicate protests, grey lines indicate hostage takings, dashed-black lines indicate shootings, and dashed-grey lines indicate bombings. The results reinforce the centrality of violent tactics and political agendas to our understanding of terrorism: protests are unlikely to be considered terrorism regardless of the motivation attributed to the actor, and attacks motivated by personal disputes are unlikely to be considered terrorism regardless of the tactics employed. Note that protests motivated by personal disputes were precluded in the experimental design as part of the randomization constraints, which is why this combination is not depicted here.

Figure 6: Interaction between the tactic and actor.



Interaction between the tactic and actor. Black lines indicate protests, grey lines indicate hostage takings, dashed-black lines indicate shootings, and dashed-grey lines indicate bombings. The results show that a history of mental illness suppresses the effects of violent tactics on the likelihood of an incident being classified as terrorism, while violent incidents carried out by more formal organizations are more likely to be classified as terrorism than those by lone individuals.

Figure 7: Interaction between the actor and their motivation.

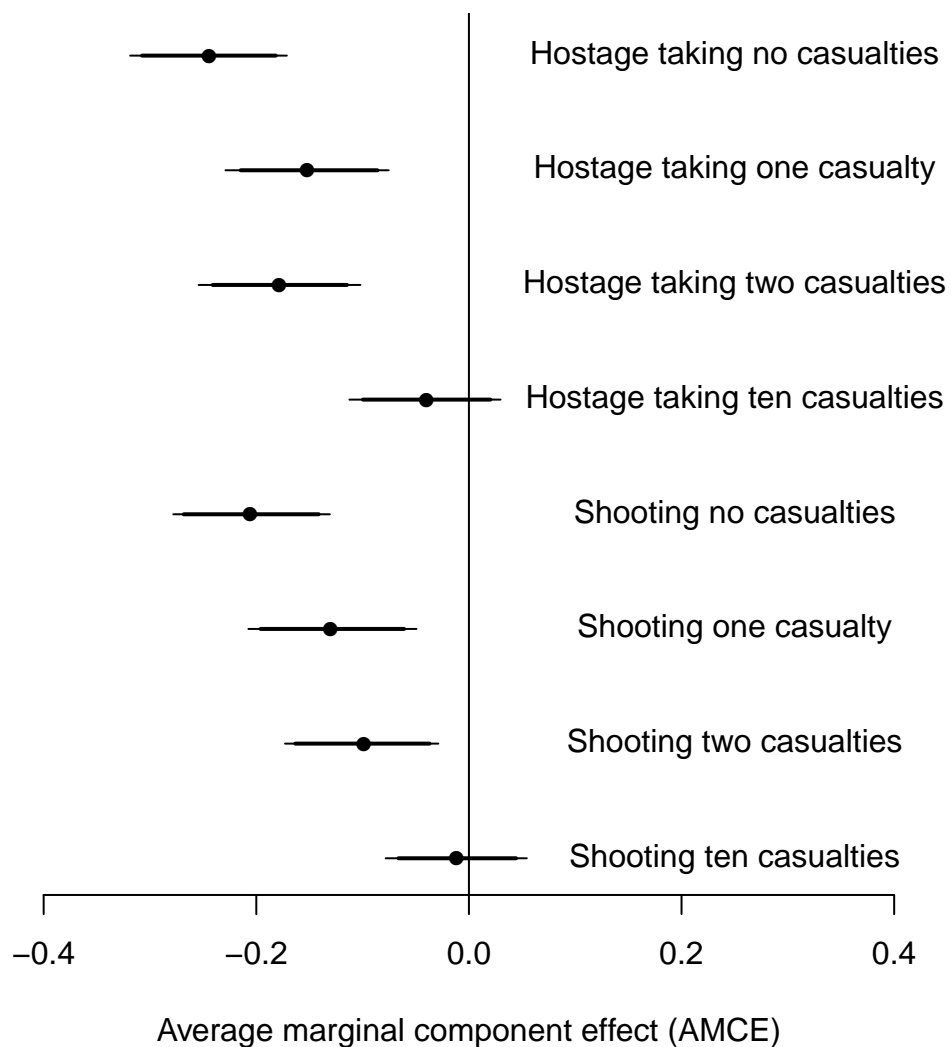


Interaction between the actor and their motivation. Black lines indicate personal disputes, grey lines indicate unclear motivations, dashed-black lines indicate hatred, dashed-grey lines indicate policy change, and dashed-black lines indicate government overthrow. Note that only collective actors were subjected to a randomization constraint that precluded them from carrying out attacks for personal disputes, which is why that combination is not depicted here.

terrorism than a hostage-taking or shooting that kills two people. Moreover, hostage-takings with ten casualties are statistically indistinguishable from bombings with zero casualties. This highlights the extent to which bombing is uniquely associated with terrorism as well as the central importance of the tactics chosen by the actor.

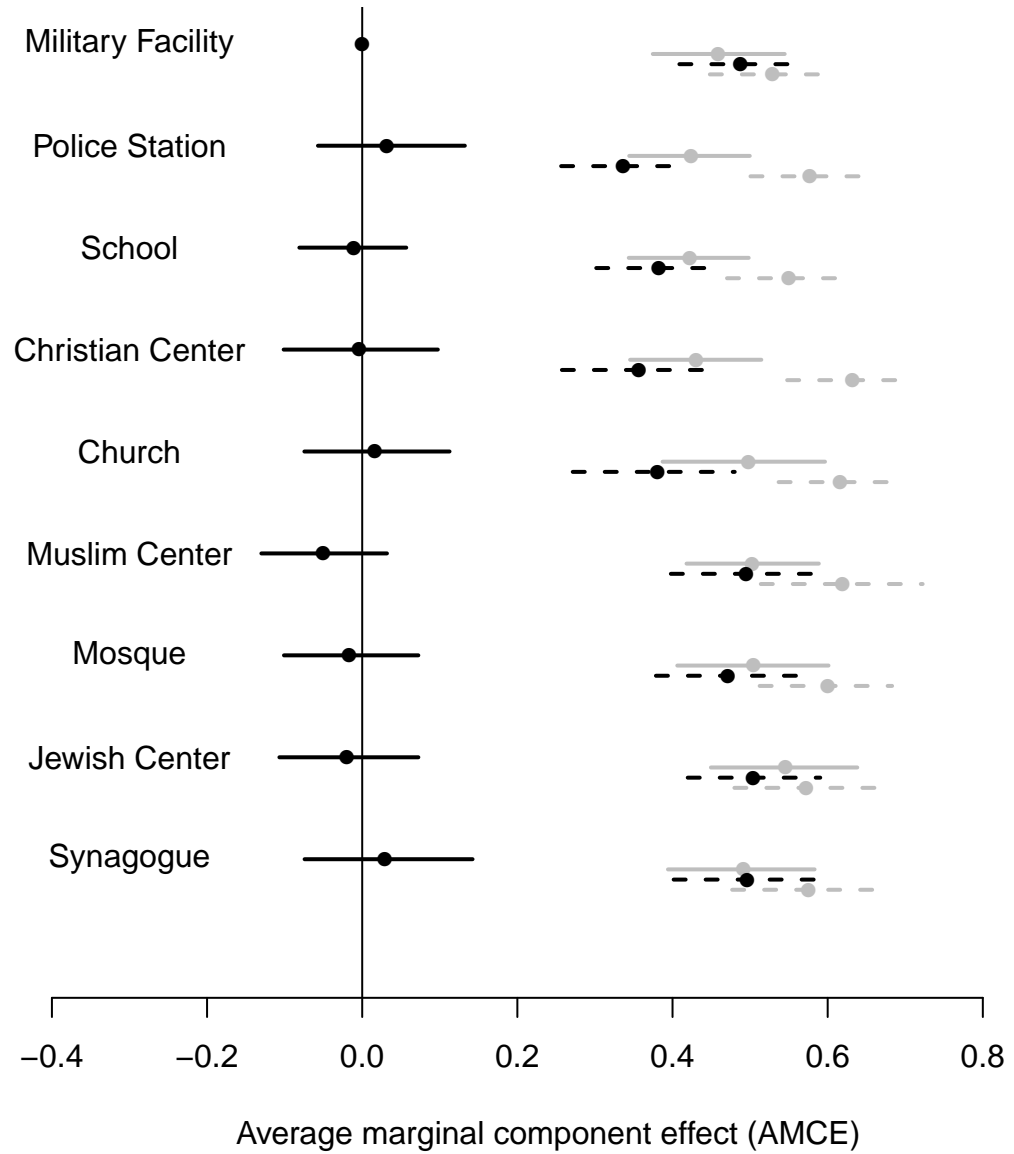
Figures 9-10 present a number of results that are of particular theoretical interest given contemporary debates as to the role that double standards against Muslims play in terrorism classifications, whether relating to Mosques and Muslim community centers being treated differently from other targets, or Muslim perpetrators being treated differently from non-Muslim ones — and, whether these effects vary between liberals and conservatives. First, Figure 9 estimates interactions between the tactic and the target. The results demonstrate that regardless of the target, violent incidents are more likely to be classified as terrorism than non-violent incidents. There also seem to be some differences between shootings and bombings for Christian centers and churches when compared to Muslim centers and mosques and Jewish centers and synagogues, but the effects are substantively small. Second, Figure 10 presents the interaction between the type of actor and their ideology divided by liberals and conservatives, demonstrating that both the direction and size of the effects are fairly similar. The absence of heterogeneous effects here is striking given debates about whether Muslim perpetrators of attacks can benefit from the same “mental illness” defense routinely given to Christian perpetrators. In short, these supplementary analyses generally fail to find evidence of much heterogeneity with respect to Muslim targets or perpetrators — consistent with the relatively small effect ranges that interactions with ideology display in Figure 3(b). See Appendix §2.3 for further explorations of the effects of anti-Muslim prejudice in our results.

Figure 8: How bombing with no casualties compares to other tactics and casualty tolls.



The AMCEs demonstrate the importance of tactics, and bombing in particular, for the likelihood an incident is defined to be terrorism. Bombing with no casualties is statistically indistinguishable from shootings and hostage takings with ten casualties.

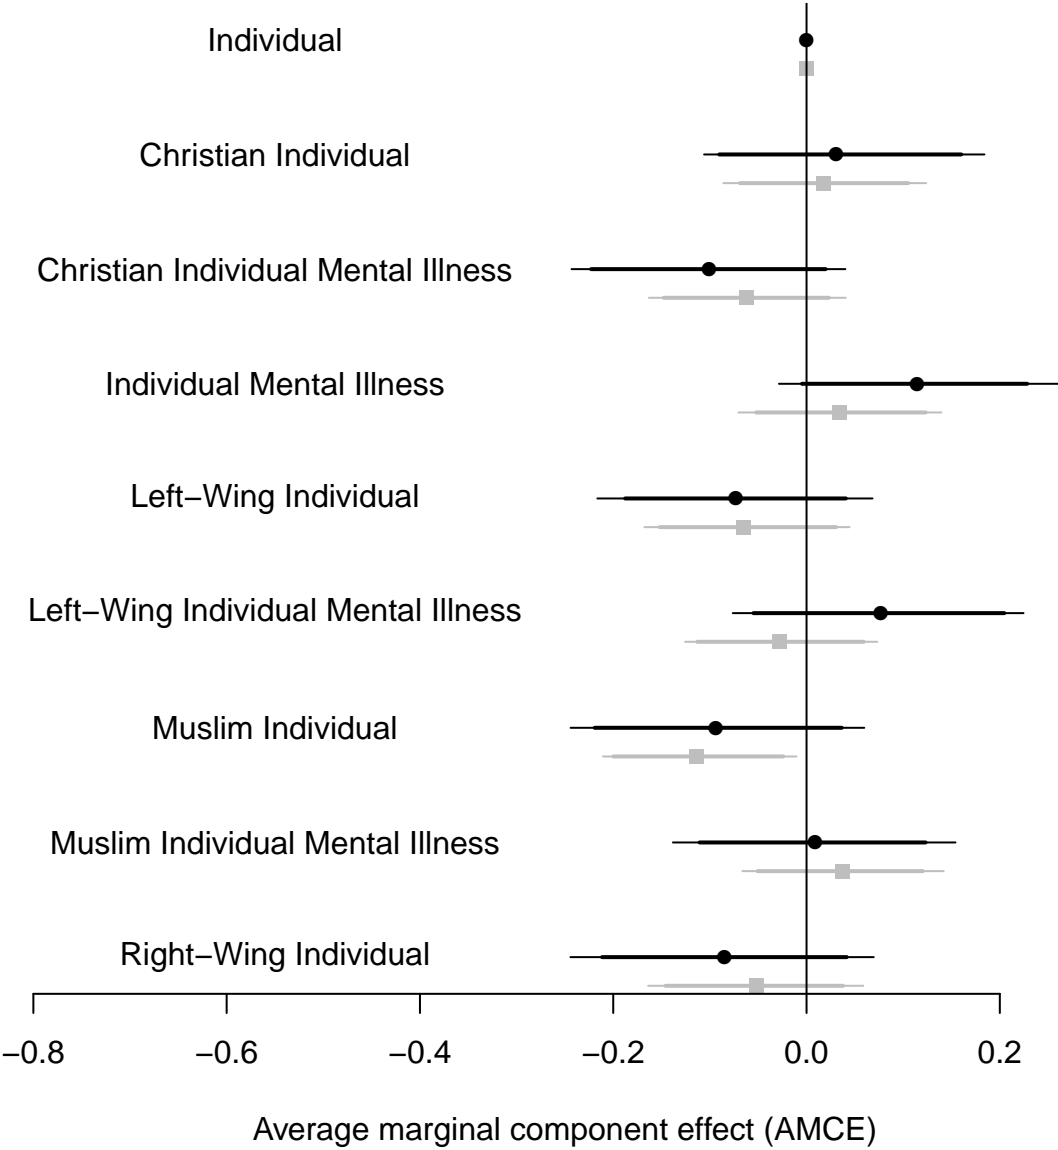
Figure 9: Interaction between the tactic and the target.



Interaction between the tactic and the target. Black lines indicate protests, grey lines indicate hostage takings, dashed-black lines indicate shootings, and dashed-grey lines indicate bombings.



Figure 10: Interaction between the actor and their identity split by liberals and conservatives.



Interaction between the type of actor and their ideology. Conservatives are denoted with black lines while liberals are denoted with grey.

## 2.3 The moderating role of attitudes towards Islam

Given recent debates about the extent to which discourses about terrorism are wrapped up in double standards against Muslims,<sup>30</sup> we explicitly included a “Muslim” treatment as part of our social categorization treatments. The results in the main text found a modest, but statistically significant effect, in which incidents by perpetrators described as Muslim were 6% more likely to be perceived as terrorism than incidents by perpetrators described as Christian, and 8% more likely to be classified as terrorism than incidents by perpetrators for whom no social categorization was provided. The effect was small, but was also larger among conservatives, whereas the effect was not statistically significant among liberals.<sup>31</sup>

We also test for heterogeneous effects in a more direct way, based on participants’ attitudes towards Islam. The survey thus borrowed instrumentation from Imhoff and Recker (2012), who, studying public attitudes towards Muslims in Germany, suggest that most measures of negative affect towards Islam actually conflate two distinct factors: “Islamoprejudice”, which refers to prejudiced views of Islam, with “Secular Critique”, a kind of liberal critique of religious practices more generally voiced by critics like Richard Dawkins or Christopher Hitchens. They find that the two constructs are conceptually distinct: Islamoprejudice is positively correlated with Right-Wing Authoritarianism (RWA) and Social Dominance Orientation (SDO), for example, while Secular Critique is negatively correlated with the former, and independent of the latter. Similarly, Islamoprejudice scores are positively correlated with Implicit Attitude Test (IAT) measures of prejudice against Muslims, while Secular Critique is not. Accordingly, in the demographic questionnaire administered to participants, we presented respondents with the five statements shown below (with standard five-item Likert response options ranging from “strongly disagree” to “strongly agree”), the first three of which measure Islamoprejudice, and latter two Secular Criticism.

1. Islam and Christianity share the same universal ethical principles.
2. Islam is an archaic religion, unable to adjust to the present.
3. I think the Islamic religion and its aggressive sides predispose it towards terrorism.

---

<sup>30</sup>See, for example, Ben Norton, “Our terrorism double standard: After Paris, let’s stop blaming Muslims and take a hard look at ourselves.” *Salon*, November 14, 2015. [http://www.salon.com/2015/11/14/our\\_terrorism\\_double\\_standard\\_after\\_paris\\_lets\\_stop\\_blaming\\_muslims\\_and\\_take\\_a\\_hard\\_look\\_at\\_ourselves/](http://www.salon.com/2015/11/14/our_terrorism_double_standard_after_paris_lets_stop_blaming_muslims_and_take_a_hard_look_at_ourselves/). Qasim Rashid, “The Terrorism Double Standard”, *Huffington Post*, December 9, 2015. [http://www.huffingtonpost.com/qasim-rashid/the-terrorism-double-standard\\_b\\_8723512.html](http://www.huffingtonpost.com/qasim-rashid/the-terrorism-double-standard_b_8723512.html). Haroon Moghul, “The only difference between Christian gunman and a Muslim terrorist is racism”, *Quartz*, March 29, 2016. <http://qz.com/649933/the-only-difference-between-a-christian-gunman-and-a-muslim-terrorist-is-racism/>. Or, conversely, Jonathan Zimmerman, “The double standard on terrorism justifications”, *New York Daily News*, March 28, 2016, <http://www.nydailynews.com/opinion/jonathan-zimmerman-double-standard-explaining-terror-article-1.2580468>.

<sup>31</sup>Because of the small magnitude of the effect, in Appendix §3.2, we test to see whether the effect shrinks over time in a manner that might be indicative of social desirability bias, but fail to find evidence of it. See also Appendix §2.2, which looks for double standards in terms of the targeting of Islamic religious facilities and community centers.


4. Religion becomes a problem when humans try to take holy writing literally.
5. The strict division between church and state is a Western accomplishment that would be progress in many Islamic countries.

In the analysis below, we mean-split each scale, and present the full set of AMCEs for each subgroup. Given the theoretical work on differentiating the two constructs from one another, we expect there to be heterogeneous treatment effects for the Islamoprejudice measure, but not the Secular Critique one. We present the full set of results in Figures 11 and 12.<sup>32</sup> Figure 11 presents heterogeneous effects by the Islamoprejudice measure. Importantly, individuals high in Islamoprejudice (in black) are more likely to perceive incidents carried out by Muslims as terrorism, and significantly less likely to perceive incidents carried by Christians as terrorism. Interestingly, they also are significantly more likely to define low-casualty incidents as terrorism, and attributed motivations appear to matter less for them than for individuals who are low in Islamoprejudice. As before, though, the magnitude of the bias is relatively small, and the overall ranking of factors remains the same. In contrast in Figure 12, individuals high in Secular Critique (in black) are no more likely to perceive attacks executed by Muslims as terrorism than individuals low in Secular Critique, although the AMCEs are positive and significant for both groups. On the whole, then, the results reconfirm the findings in the main text: while incidents carried out by Muslim perpetrators are indeed more likely to be characterized as terrorism, and the treatment effect is largest amongst subgroups with negative views of Islam, the effects on the whole are modest in size, and the most prejudiced respondents nonetheless appear to define events as terrorism using a similar ranking of indicators as their low-prejudice counterparts.

---


<sup>32</sup>We present the complete set of results rather than just the effect of the Muslim treatment, because of the possibility of prejudice spilling over into the other treatments: for example, it is possible that an individual who tends to equate Islam with terrorism would be suspicious of Muslim protests, or perceive stronger treatment effects for the unclear motive treatment, automatically attributing the attack to radical Islamism.

Figure 11: Average Marginal Component Effects by levels of Islamoprejudice



Black lines denote individuals with high levels of Islamoprejudice while grey lines denote those with low levels. Results demonstrate that individuals with high levels of Islamoprejudice tend to be more likely to classify incidents carried out by Muslims as terrorism, but the effect size is modest, and the overall ranking of factors similar to those who have low levels of prejudice against Islam.

Figure 12: Average Marginal Component Effects by levels of Secular Critique against Islam



Black lines denote individuals with high levels of a secular critique of Islam while grey lines denote those with low levels. Results demonstrate that individuals with high and low levels of making a secular critique of Islam generally perceive the same factors as important in shaping their perception of whether incidents should be classified as terrorism.

### 3 Model diagnostics

#### 3.1 Randomization checks

Randomization checks for the seven different treatment categories. Results demonstrate that there is balance across treatment categories across respondent covariates which we used to generate the weighted estimates presented in the main text.

	Hostage Taking	Shooting	Bombing
Intercept	[-2.576, 17.417]	[-2.861, 16.448]	[0.139, 19.442]
Male	[-0.082, 0.137]	[-0.040, 0.182]	[-0.131, 0.084]
Birthyear	[-0.009, 0.002]	[-0.008, 0.002]	[-0.010, 0.000]
Education	[-0.031, 0.055]	[-0.033, 0.055]	[-0.047, 0.033]
Ideology	[-0.036, 0.036]	[-0.042, 0.028]	[-0.035, 0.035]

Table 2: Tactic Randomization.

	One Casualty	Two Casualties	Ten Casualties
Intercept	[-3.129, 15.002]	[-6.211, 9.716]	[-3.058, 14.770]
Male	[0.000, 0.249]	[-0.126, 0.069]	[-0.031, 0.167]
Birthyear	[-0.008, 0.001]	[-0.005, 0.003]	[-0.008, 0.001]
Education	[-0.082, 0.000]	[-0.027, 0.051]	[-0.019, 0.054]
Ideology	[-0.040, 0.019]	[-0.031, 0.029]	[-0.030, 0.033]

Table 3: Casualty Randomization.

	US Ties	Foreign Ties	Group	Individual	Individual Mental
Intercept	[-15.429, 13.857]	[-25.542, 0.000]	[-19.500, 2.793]	[-10.866, 9.882]	[-21.476, 0.000]
Male	[-0.095, 0.249]	[-0.098, 0.160]	[-0.095, 0.165]	[-0.074, 0.173]	[-0.073, 0.163]
Birthyear	[-0.007, 0.007]	[0.000, 0.013]	[-0.001, 0.010]	[-0.005, 0.006]	[0.000, 0.011]
Education	[-0.106, 0.023]	[-0.091, 0.000]	[-0.072, 0.025]	[-0.084, 0.007]	[-0.056, 0.034]
Ideology	[-0.049, 0.048]	[-0.034, 0.036]	[-0.012, 0.065]	[-0.023, 0.045]	[-0.035, 0.034]

Table 4: Actor Randomization.

	Christian	Muslim	Left-Wing	Right-Wing
Intercept	[-7.328, 13.244]	[-3.425, 15.550]	[-6.692, 12.341]	[0.000, 19.213]
Male	[-0.053, 0.164]	[-0.030, 0.202]	[0.000, 0.229]	[-0.027, 0.184]
Birthyear	[-0.007, 0.004]	[-0.008, 0.002]	[-0.006, 0.003]	[-0.010, 0.000]
Education	[-0.030, 0.056]	[-0.035, 0.046]	[-0.023, 0.062]	[-0.030, 0.051]
Ideology	[-0.060, 0.005]	[-0.035, 0.046]	[-0.037, 0.031]	[-0.031, 0.031]

Table 5: Identity Randomization.

	Gov Overthrow	Policy Change	Hatred	Personal Dispute
Intercept	[-6.612, 14.248]	[-5.165, 12.837]	[-5.283, 13.338]	[-3.500, 16.466]
Male	[-0.138, 0.114]	[-0.036, 0.186]	[-0.120, 0.105]	[-0.132, 0.116]
Birthyear	[-0.007, 0.003]	[-0.007, 0.003]	[-0.007, 0.003]	[-0.008, 0.002]
Education	[-0.040, 0.054]	[-0.031, 0.052]	[-0.036, 0.048]	[-0.035, 0.060]
Ideology	[-0.049, 0.024]	[-0.040, 0.025]	[-0.064, 0.001]	[-0.067, 0.008]

Table 6: Motivation Randomization.

	Police Station	School	Christian Center	Muslim Center	Jewish Center	Church	Mosque	Synagogue
Intercept	[-6.125, 9.434]	[-7.277, 9.491]	[-6.691, 11.254]	[-21.234, 0.000]	[-12.389, 7.693]	[-22.897, 0.000]	[-5.913, 14.008]	[-11.170, 7.653]
Male	[-0.141, 0.047]	[-0.087, 0.115]	[-0.168, 0.062]	[-0.133, 0.073]	[0.000, 0.306]	[-0.180, 0.073]	[-0.151, 0.095]	[-0.085, 0.134]
Birthyear	[-0.005, 0.003]	[-0.005, 0.004]	[-0.006, 0.003]	[0.000, 0.010]	[-0.004, 0.006]	[-0.001, 0.011]	[-0.007, 0.003]	[-0.004, 0.005]
Education	[-0.032, 0.044]	[-0.009, 0.073]	[-0.018, 0.077]	[-0.033, 0.053]	[-0.042, 0.050]	[0.000, 0.117]	[-0.046, 0.046]	[-0.023, 0.066]
Ideology	[-0.010, 0.070]	[-0.035, 0.049]	[-0.061, 0.020]	[-0.053, 0.021]	[-0.035, 0.042]	[-0.012, 0.061]	[-0.065, 0.018]	[-0.074, 0.009]

Table 7: Target Randomization.

	For Dem	For Dem HR	For Dict	For Dict HR
Intercept	[-4.281, 14.670]	[-3.403, 16.130]	[-7.546, 11.185]	[-3.847, 14.556]
Male	[-0.131, 0.097]	[-0.154, 0.070]	[-0.196, 0.030]	[-0.175, 0.050]
Birthyear	[-0.008, 0.002]	[-0.009, 0.001]	[-0.006, 0.004]	[-0.008, 0.002]
Education	[-0.045, 0.035]	[-0.072, 0.011]	[-0.050, 0.032]	[-0.068, 0.016]
Ideology	[-0.018, 0.043]	[0.000, 0.078]	[-0.026, 0.039]	[-0.013, 0.045]

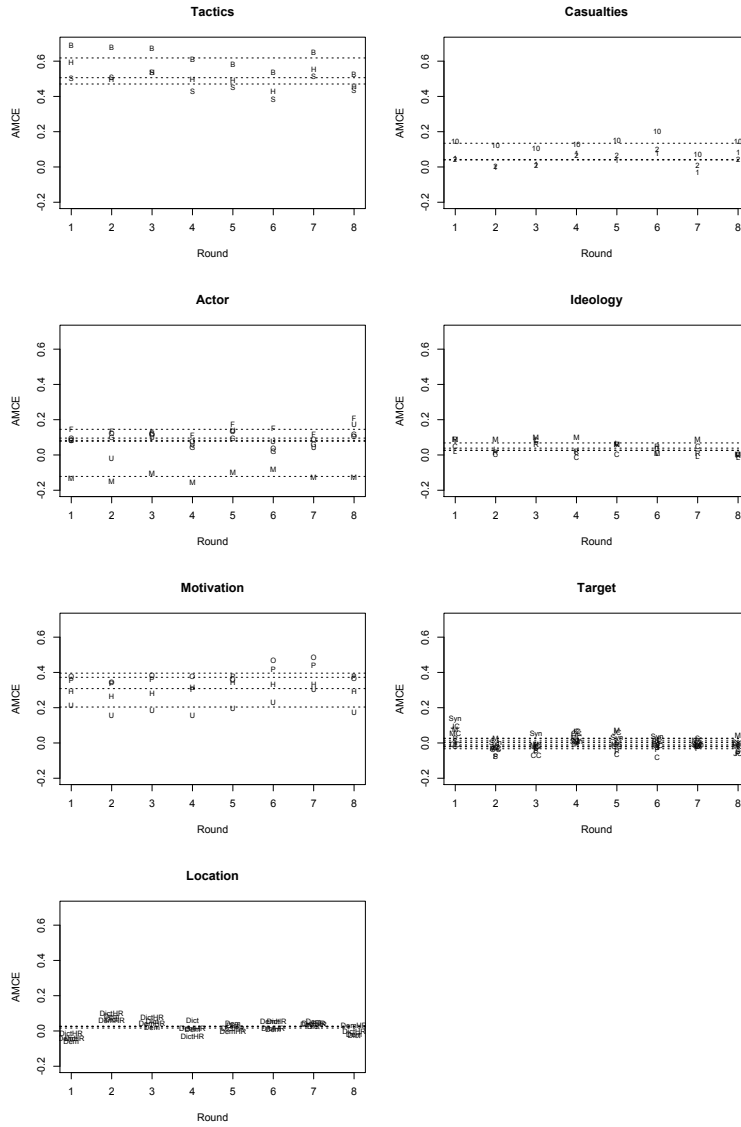
Table 8: Country Location Randomization.

### 3.2 Carryover assumption test

One of the assumptions underlying conjoint analyses is that the results do not display over-time effects, thereby enabling us to pool results across multiple rounds. Figure 13 presents AMCEs within each round, with each letter depicting a different level of the treatment (e.g. for tactics, “H” = hostage taking, “S” = shooting, and “B” = bombing), and the dashed horizontal lines depicting the AMCEs averaged across all eight rounds. The figure shows the results to be relatively stable over time: for the Motivation factor, for example, unclear motivations (depicted by the “U” in the plot) tend to hover around the cross-round AMCE of 0.2; overthrowing the government (“O”) and policy change (“P”) tend to have the largest effects across all rounds, and so on.

The absence of longitudinal effects is also substantively interesting in the case of ideology. One interpretation of the relatively modest effects for Muslim actors is social desirability bias: that respondents underreport the extent to which they characterize incidents with Muslim perpetrators, because they are aware of the purpose of the study. If this were the case, one empirical implication would be that the effect of the Muslim treatment should shrink over time: respondents who receive the Muslim treatment in the first round, for example, will not have observed the other levels of the treatment, but have an 80% probability of observing a non-Muslim actor in each subsequent round, such that the potential for demand effects would increase. Yet the ideology panel shows that the effect of the Muslim treatment remains relatively constant, thereby failing to provide evidence in support of a social desirability explanation.

Figure 13: AMCEs display little evidence of longitudinal effects.





## 4 Alternative predicted probabilities

Tables 9-11 present a number of alternative estimates of the probability a number of incidents are classified as terrorism. The results generally demonstrate consistency in the types of incidents we should expect to be most likely to be classified as terrorism, regardless of the statistical model used. Table 9 presents predicted probabilities generated by a Support Vector Machine classifier with LASSO constraints with no weights. Predicted probabilities generated by a Linear Probability Model without weights are presented in Table 10 while Table 11 presents probabilities from a Linear Probability Model with weights included. Note that when using a Linear Probability Model there is a larger range of plausible values than the predicted probabilities presented in the paper, since the predicted values from the SVM are calculated using a logistic transformation. However, the rank ordering among incidents is largely the same.

Table 9: Predicted probabilities: Support Vector Machine (SVM) classifier

Incident	Date	Predicted Probability
Dallas Police HQ Shooting	(06/13/15)	0.29
University of California Tuition Hike Protests	(03/18/15)	0.32
UCLA Black Lives Matter Protest	(10/08/15)	0.33
East Selma Church Shooting	(09/20/15)	0.33
Islamic Community Center of Phoenix Demonstrations	(10/10/15)	0.34
Newport Church hostage situation	(07/30/06)	0.35
Shooting of Police in Oakland	(03/21/09)	0.41
Marysville Pilchuck High School shooting	(10/24/14)	0.42
Rocori High School shooting	(09/24/03)	0.42
Rosemary Anderson High School shooting	(12/13/14)	0.44
Camp Shelby Shootings	(08/05/15)	0.45
St. Columbanus Church Shooting	(11/26/12)	0.45
Zvornik Police Station shooting	(04/27/15)	0.48
Poe Elementary School Bombing	(09/15/59)	0.49
Shooting of Police in Brooklyn	(12/20/14)	0.49
University of Alabama Huntsville	(12/02/10)	0.49
Charleston church Shooting	(06/17/15)	0.50
Pentagon Metro Shooting	(03/04/10)	0.50
Knoxville Unitarian Universalist Church shooting	(07/27/08)	0.53
Overland Park Jewish Community Center Shooting	(04/13/14)	0.54
Seattle Jewish Federation Shooting	(07/28/06)	0.55
Copenhagen Synagogue Shooting	(02/14/15)	0.56
Shooting of George Tiller	(05/31/09)	0.59
Nag Hammadi massacre	(01/07/10)	0.61
Lombard Islamic School bombing	(08/12/12)	0.63
Contra attack in Quilali	(11/11/87)	0.65
Chattanooga shootings	(07/16/15)	0.67
Fort Hood Shootings	(11/05/09)	0.67
ETA Sanguesa car bombing	(05/30/03)	0.68
Bombing of Shiraa village mosque	(12/30/14)	0.69
Kehilat Bnei Torah synagogue attack	(11/18/14)	0.69
Porte de Vincennes hostage situation	(01/09/15)	0.69
Zif School Bombing	(09/17/02)	0.70
KKK Selma Bombing	(9/15/63)	0.70
Hamas attack on IDF in Khan Yunis	(12/24/14)	0.70
Aksu bombing	(08/19/10)	0.72
Pakistan Army General HQ hostage situation	(10/10/09)	0.72
Camp Integrity Suicide bombing	(08/07/15)	0.74
Shebaa Farms incident	(01/28/15)	0.76

Table 10: Predicted probabilities: Linear Probability Model (LPM)

Incident	Date	Predicted Probability
University of California Tuition Hike Protests	(03/18/15)	0.18
Dallas Police HQ Shooting	(06/13/15)	0.19
UCLA Black Lives Matter Protest	(10/08/15)	0.21
Newport Church hostage situation	(07/30/06)	0.24
East Selma Church Shooting	(09/20/15)	0.24
Islamic Community Center of Phoenix Demonstrations	(10/10/15)	0.26
Rosemary Anderson High School shooting	(12/13/14)	0.32
St. Columbanus Church Shooting	(11/26/12)	0.34
Camp Shelby Shootings	(08/05/15)	0.42
Shooting of Police in Oakland	(03/21/09)	0.43
Marysville Pilchuck High School shooting	(10/24/14)	0.45
Rocori High School shooting	(09/24/03)	0.45
Pentagon Metro Shooting	(03/04/10)	0.51
Zvornik Police Station shooting	(04/27/15)	0.52
University of Alabama Huntsville	(12/02/10)	0.54
Charleston church Shooting	(06/17/15)	0.56
Shooting of Police in Brooklyn	(12/20/14)	0.57
Poe Elementary School Bombing	(09/15/59)	0.57
Knoxville Unitarian Universalist Church shooting	(07/27/08)	0.59
Seattle Jewish Federation Shooting	(07/28/06)	0.59
Overland Park Jewish Community Center Shooting	(04/13/14)	0.62
Shooting of George Tiller	(05/31/09)	0.65
Lombard Islamic School bombing	(08/12/12)	0.68
Nag Hammadi massacre	(01/07/10)	0.74
Chattanooga shootings	(07/16/15)	0.80
Fort Hood Shootings	(11/05/09)	0.80
Contra attack in Quilali	(11/11/87)	0.81
Kehilat Bnei Torah synagogue attack	(11/18/14)	0.83
Porte de Vincennes hostage situation	(01/09/15)	0.86
ETA Sanguesa car bombing	(05/30/03)	0.89
Bombing of Shiraa village mosque	(12/30/14)	0.89
KKK Selma Bombing	(9/15/63)	0.90
Hamas attack on IDF in Khan Yunis	(12/24/14)	0.90
Zif School Bombing	(09/17/02)	0.90
Pakistan Army General HQ hostage situation	(10/10/09)	0.98
Aksu bombing	(08/19/10)	1.00
Copenhagen Synagogue Shooting	(02/14/15)	1.00
Camp Integrity Suicide bombing	(08/07/15)	1.00
Shebaa Farms incident	(01/28/15)	1.00

Table 11: Predicted probabilities: Linear Probability Model (LPM) with weights

Incident	Date	Predicted Probability
Dallas Police HQ Shooting	(06/13/15)	0.16
University of California Tuition Hike Protests	(03/18/15)	0.18
UCLA Black Lives Matter Protest	(10/08/15)	0.22
Newport Church hostage situation	(07/30/06)	0.23
East Selma Church Shooting	(09/20/15)	0.24
Islamic Community Center of Phoenix Demonstrations	(10/10/15)	0.30
Rosemary Anderson High School shooting	(12/13/14)	0.32
St. Columbanus Church Shooting	(11/26/12)	0.35
Camp Shelby Shootings	(08/05/15)	0.39
Shooting of Police in Oakland	(03/21/09)	0.41
Marysville Pilchuck High School shooting	(10/24/14)	0.41
Rocori High School shooting	(09/24/03)	0.41
Pentagon Metro Shooting	(03/04/10)	0.45
Zvornik Police Station shooting	(04/27/15)	0.49
University of Alabama Huntsville	(12/02/10)	0.50
Poe Elementary School Bombing	(09/15/59)	0.53
Shooting of Police in Brooklyn	(12/20/14)	0.54
Seattle Jewish Federation Shooting	(07/28/06)	0.57
Charleston church Shooting	(06/17/15)	0.57
Knoxville Unitarian Universalist Church shooting	(07/27/08)	0.60
Overland Park Jewish Community Center Shooting	(04/13/14)	0.63
Shooting of George Tiller	(05/31/09)	0.65
Lombard Islamic School bombing	(08/12/12)	0.65
Nag Hammadi massacre	(01/07/10)	0.74
Chattanooga shootings	(07/16/15)	0.76
Fort Hood Shootings	(11/05/09)	0.76
Kehilat Bnei Torah synagogue attack	(11/18/14)	0.80
Contra attack in Quilali	(11/11/87)	0.80
Porte de Vincennes hostage situation	(01/09/15)	0.82
ETA Sanguesa car bombing	(05/30/03)	0.84
Hamas attack on IDF in Khan Yunis	(12/24/14)	0.87
Bombing of Shiraa village mosque	(12/30/14)	0.88
Zif School Bombing	(09/17/02)	0.89
KKK Selma Bombing	(9/15/63)	0.92
Pakistan Army General HQ hostage situation	(10/10/09)	0.93
Aksu bombing	(08/19/10)	0.99
Copenhagen Synagogue Shooting	(02/14/15)	0.99
Camp Integrity Suicide bombing	(08/07/15)	1.00
Shebaa Farms incident	(01/28/15)	1.00

## 5 Additional information about the media content analysis

We searched LexisNexis for all articles within two weeks of the shootings in Fort Hood and Charleston that discussed these incidents. The search terms used allowed for articles that either mentioned the name of the shooter or the location of the incident and either the words massacre or shooting. The full list of search terms for Charleston was: ““Dylann Roof” OR “Charleston Shooting” OR “Charleston Church Shooting” OR “Charleston shooting” OR “Charleston church shooting” OR “Charleston Massacre” OR “Charleston Church Massacre” OR “Charleston massacre” OR “Charleston church massacre”.” The full list of search terms for Fort Hood was: “Nidal Hasan” OR “Fort Hood Shooting” OR “Fort Hood shooting” OR “Fort Hood Massacre” OR “Fort Hood massacre”.” These searches yielded a total of 1093 articles for the shooting in Charleston and 645 for the shooting in Fort Hood. In order to remove duplicate articles, the searches automatically removed highly similar articles, although the results are substantively similar regardless of whether this search restriction is used.